

Simple Linear Regression


Extra



Dinh-Thang Duong – TA
Anh-Khoi Nguyen – STA

Getting Started

❖ Objectives



Our objectives:

- Introduction to Linear Regression.
- Discuss the idea of how Linear Regression learn to solve a Regression problem.
- Discuss about Gradient Descent.
- Implement a simple version of Linear Regression using Gradient Descent.


Outline

- Introduction
- Simple Linear Regression
- Code Implementation
- Question

Introduction

Introduction

❖ Getting Started



Machine Learning (ML): A branch of AI and Computer Science which focuses on the use of data and algorithms to imitate the way that humans learn, gradually improving its accuracy.

Introduction

❖ Getting Started


Suppose you got some dataset:

	Name	Team	Number	Position	Age	Height	Weight	College	Salary
0	Avery Bradley	Boston Celtics	0.0	PG	25.0	6-2	180.0	Texas	7730337.0
1	Jae Crowder	Boston Celtics	99.0	SF	25.0	6-6	235.0	Marquette	6796117.0
2	John Holland	Boston Celtics	30.0	SG	27.0	6-5	205.0	Boston University	NaN
3	R.J. Hunter	Boston Celtics	28.0	SG	22.0	6-5	185.0	Georgia State	1148640.0
4	Jonas Jerebko	Boston Celtics	8.0	PF	29.0	6-10	231.0	NaN	5000000.0
5	Amir Johnson	Boston Celtics	90.0	PF	29.0	6-9	240.0	NaN	12000000.0
6	Jordan Mickey	Boston Celtics	55.0	PF	21.0	6-8	235.0	LSU	1170960.0
7	Kelly Olynyk	Boston Celtics	41.0	C	25.0	7-0	238.0	Gonzaga	2165160.0
8	Terry Rozier	Boston Celtics	12.0	PG	22.0	6-2	190.0	Louisville	1824360.0
9	Marcus Smart	Boston Celtics	36.0	PG	22.0	6-4	220.0	Oklahoma State	3431040.0
10	Jared Sullinger	Boston Celtics	7.0	C	24.0	6-9	260.0	Ohio State	2569260.0
11	Isaiah Thomas	Boston Celtics	4.0	PG	27.0	5-9	185.0	Washington	6912869.0
12	Evan Turner	Boston Celtics	11.0	SG	27.0	6-7	220.0	Ohio State	3425510.0
13	James Young	Boston Celtics	13.0	SG	20.0	6-6	215.0	Kentucky	1749840.0

And you want to make a program to automatically predict value of 1 column based on others.

Input

Name,
Team,
Number,
Position,
Age,
Height,
Weight,
College



Output

Salary

Introduction


❖ Getting Started

Input X

Output Y

	Name	Team	Number	Position	Age	Height	Weight	College	Salary
0	Avery Bradley	Boston Celtics	0.0	PG	25.0	6-2	180.0	Texas	7730337.0
1	Jae Crowder	Boston Celtics	99.0	SF	25.0	6-6	235.0	Marquette	6796117.0
2	John Holland	Boston Celtics	30.0	SG	27.0	6-5	205.0	Boston University	NaN
3	R.J. Hunter	Boston Celtics	28.0	SG	22.0	6-5	185.0	Georgia State	1148640.0
4	Jonas Jerebko	Boston Celtics	8.0	PF	29.0	6-10	231.0	NaN	5000000.0
5	Amir Johnson	Boston Celtics	90.0	PF	29.0	6-9	240.0	NaN	2000000.0
6	Jordan Mickey	Boston Celtics	55.0	PF	21.0	6-8	235.0	LSU	1170960.0
7	Kelly Olynyk	Boston Celtics	41.0	C	25.0	7-0	238.0	Gonzaga	2165160.0
8	Terry Rozier	Boston Celtics	12.0	PG	22.0	6-2	190.0	Louisville	1824360.0
9	Marcus Smart	Boston Celtics	36.0	PG	22.0	6-4	220.0	Oklahoma State	3431040.0
10	Jared Sullinger	Boston Celtics	7.0	C	24.0	6-9	260.0	Ohio State	2569260.0
11	Isaiah Thomas	Boston Celtics	4.0	PG	27.0	5-9	185.0	Washington	6912869.0
12	Evan Turner	Boston Celtics	11.0	SG	27.0	6-7	220.0	Ohio State	3425510.0
13	James Young	Boston Celtics	13.0	SG	20.0	6-6	215.0	Kentucky	1749840.0

Using this data to
“train” an ML
Model



Input

Name,
Team,
Number,
Position,
Age,
Height,
Weight,
College



Output

Salary

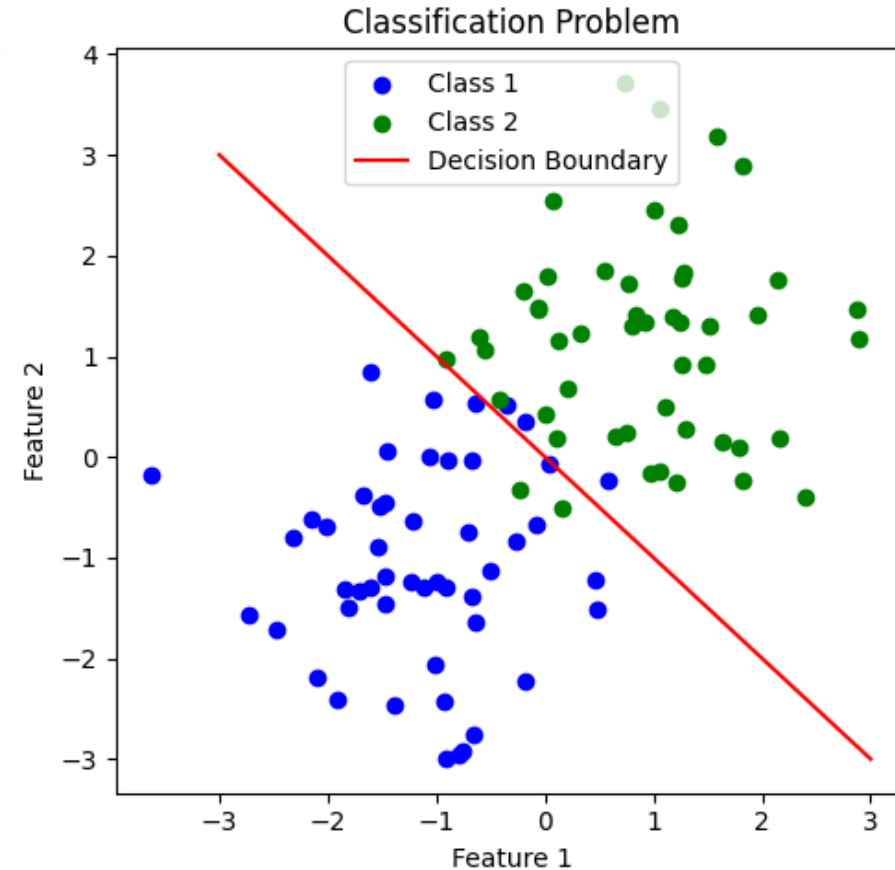
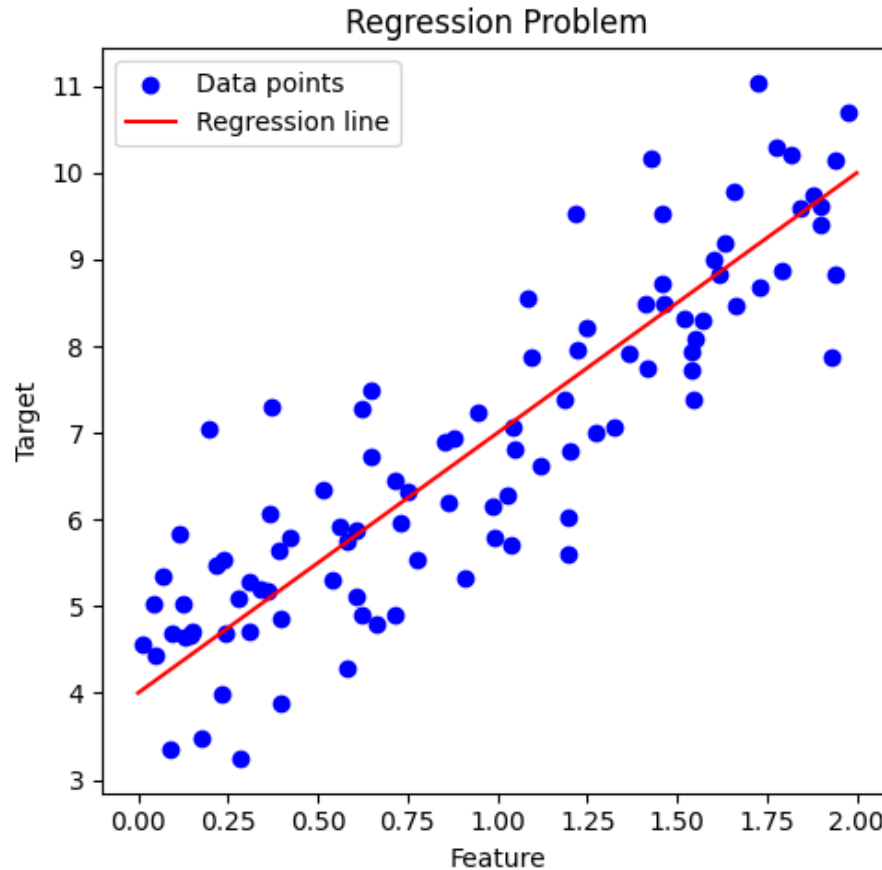
Use the trained model to predict
Salary based on any given input

Since we **use a labeled dataset to train** ML Model.

=> This is called **Supervised-learning**.

Introduction

❖ Supervised Learning


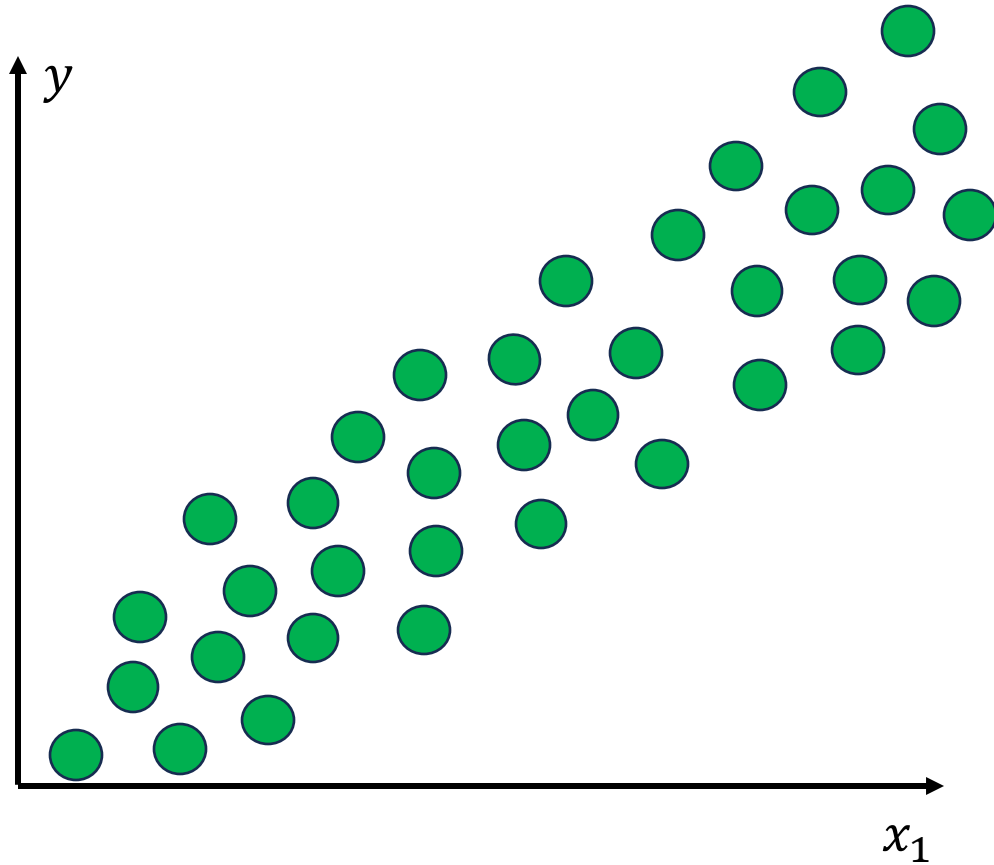


In ML Supervised-learning algorithms, we often deal with Regression and Classification

Introduction

❖ Supervised Learning: Regression

Regression: A task involving predicting a **continuous value** based on given inputs.




In general, we want to find the line that best fit the data distribution.

Simple Linear Regression

Simple Linear Regression


❖ Intro to Linear Regression



Experience	Salary
3	60
4	55
5	66
6	93
7	?

Simple Linear Regression

❖ Intro to Linear Regression

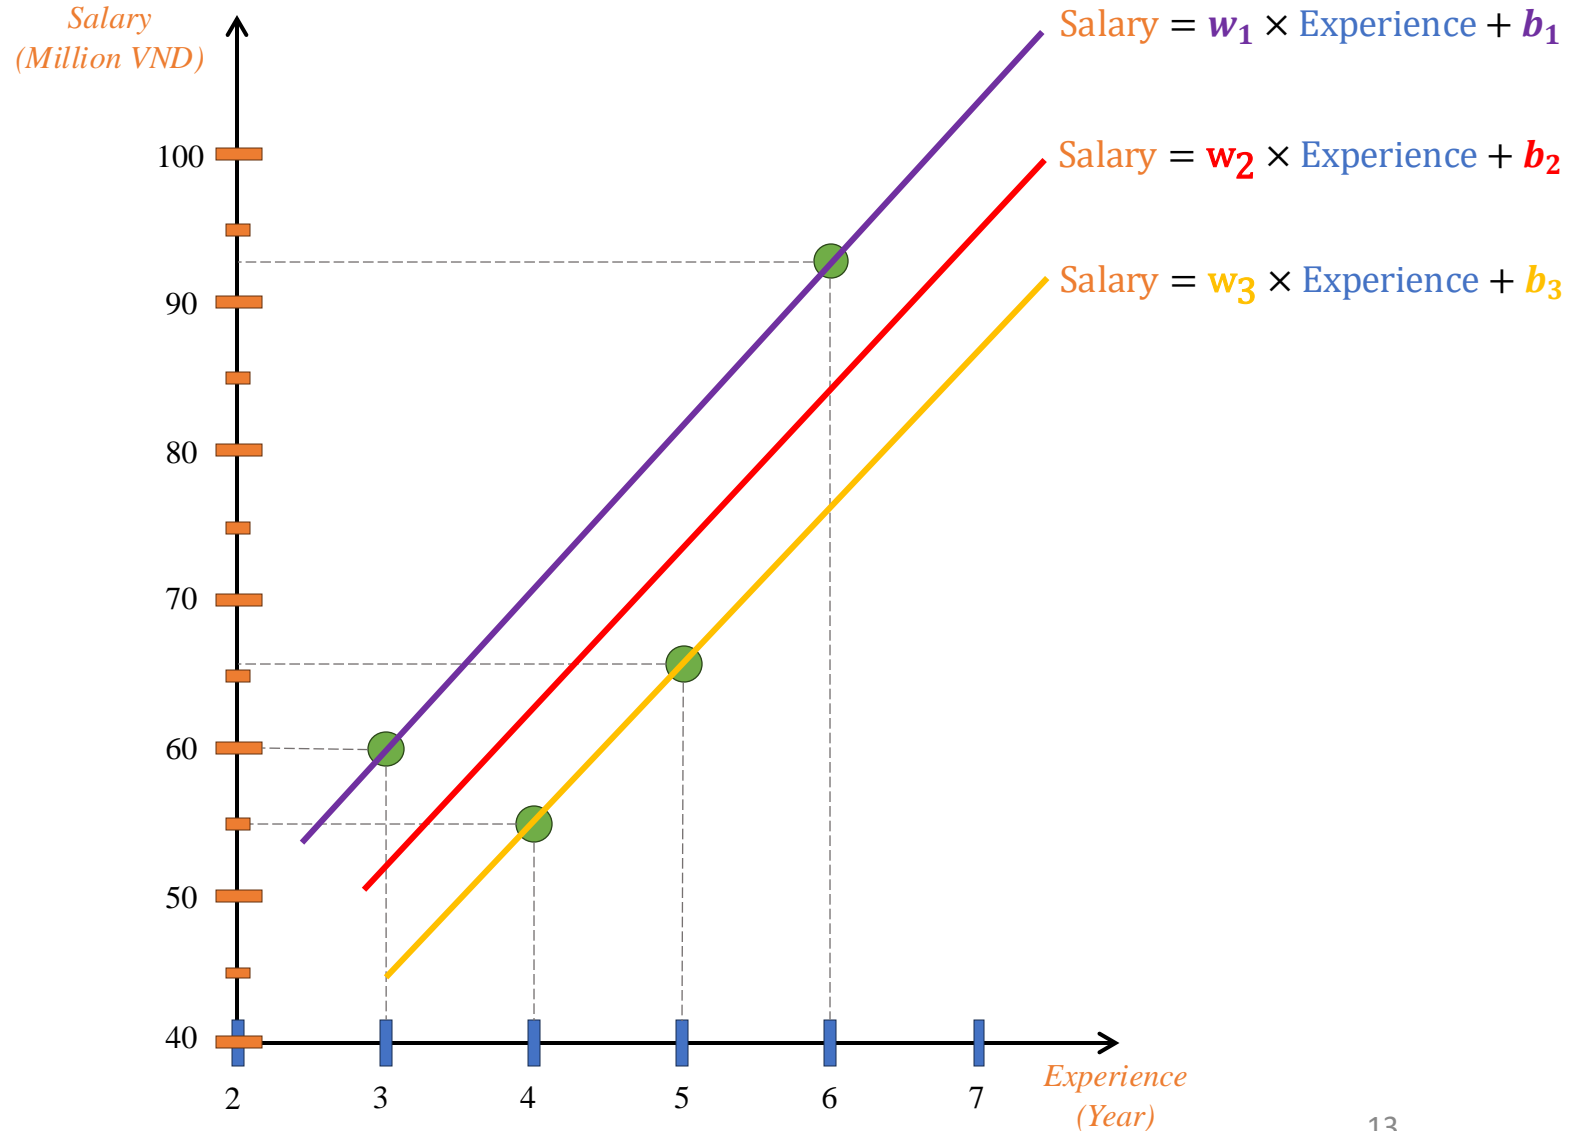


Experience	Salary
3	60
4	55
5	66
6	93
7	?

Simple Linear Regression

❖ Intro to Linear Regression


Experience	Salary
3	60
4	55
5	66
6	93
7	?



Simple Linear Regression

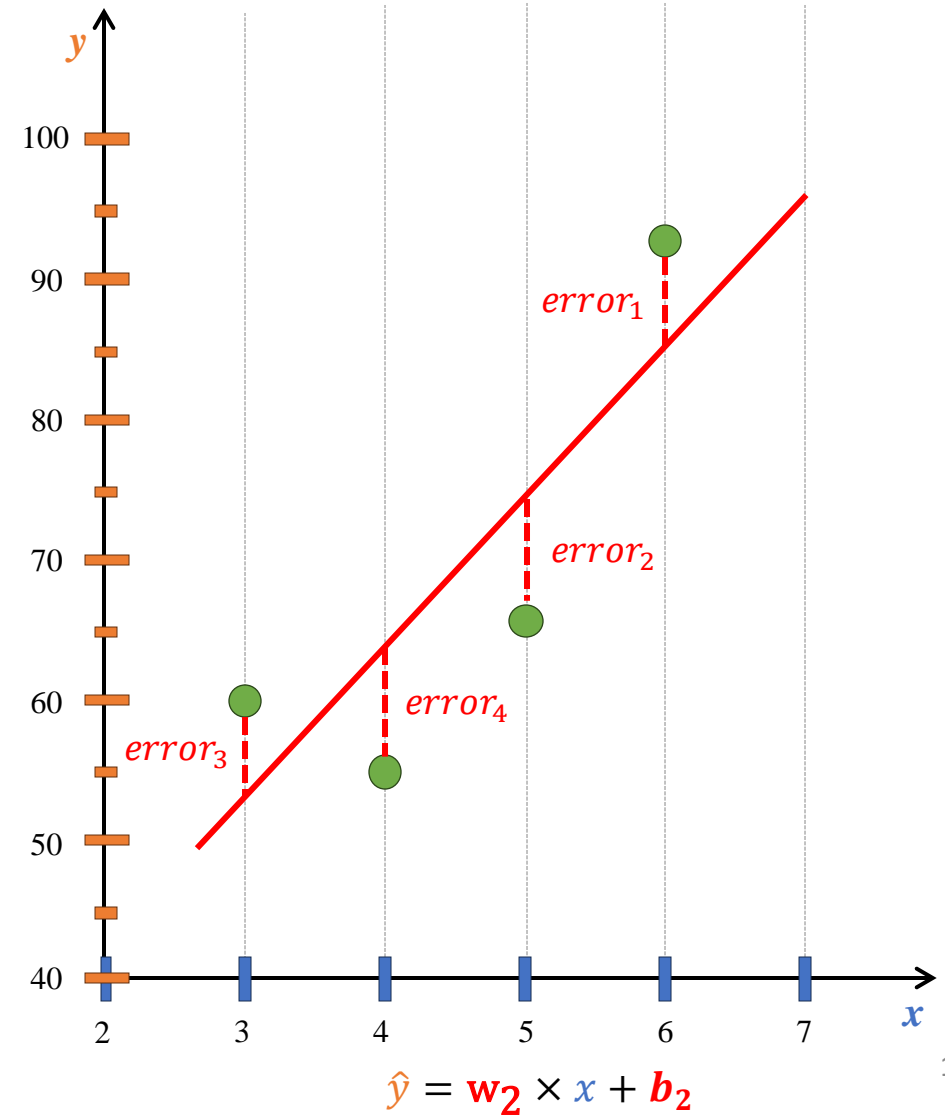

❖ Intro to Linear Regression

Experience	Salary
3	60
4	55
5	66
6	93
7	?




Simple Linear Regression

❖ Intro to Linear Regression



Simple Linear Regression



$$\text{predicted_salary} = w \times \text{Experience} + b$$

$$\text{error} = (\text{predicted_salary} - \text{real_salary})^2$$

$$\hat{y}_i = wx_i + b$$


$$\text{error}(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

$$w = 10$$

$$b = 5$$

	Experience	Salary	Predicted	Error
1	3	60	$wx_i + b$	$(\hat{y}_i - y_i)^2$
2	4	55	$wx_i + b$	$(\hat{y}_i - y_i)^2$
3	5	66	$wx_i + b$	$(\hat{y}_i - y_i)^2$
4	6	93	$wx_i + b$	$(\hat{y}_i - y_i)^2$

Simple Linear Regression



$$\text{predicted_salary} = w \times \text{Experience} + b$$

$$\text{error} = (\text{predicted_salary} - \text{real_salary})^2$$

$$\hat{y}_i = wx_i + b$$


$$\text{error}(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

$$w = 10$$

$$b = 5$$

	Experience	Salary	Predicted	Error
1	3	60	$10 \cdot 3 + 5 = 35$	
2	4	55		
3	5	66		
4	6	93		

Simple Linear Regression



$$\text{predicted_salary} = w \times \text{Experience} + b$$

$$\text{error} = (\text{predicted_salary} - \text{real_salary})^2$$

$$\hat{y}_i = wx_i + b$$


$$\text{error}(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

$$w = 10$$

$$b = 5$$

	Experience	Salary	Predicted	Error
1	3	60	35	$(60 - 35)^2 = 625$
2	4	55		
3	5	66		
4	6	93		

Simple Linear Regression



$$\text{predicted_salary} = w \times \text{Experience} + b$$

$$\text{error} = (\text{predicted_salary} - \text{real_salary})^2$$


$$\hat{y}_i = wx_i + b$$

$$\text{error}(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

$$w = 10$$
$$b = 5$$

	Experience	Salary	Predicted	Error
1	3	60	35	625
2	4	55	$10 \cdot 4 + 5 = 45$	$(55 - 45)^2 = 100$
3	5	66	$10 \cdot 5 + 5 = 55$	$(66 - 55)^2 = 121$
4	6	93	$10 \cdot 6 + 5 = 65$	$(93 - 65)^2 = 784$

Simple Linear Regression



$$\hat{y}_i = wx_i + b$$

$$error(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

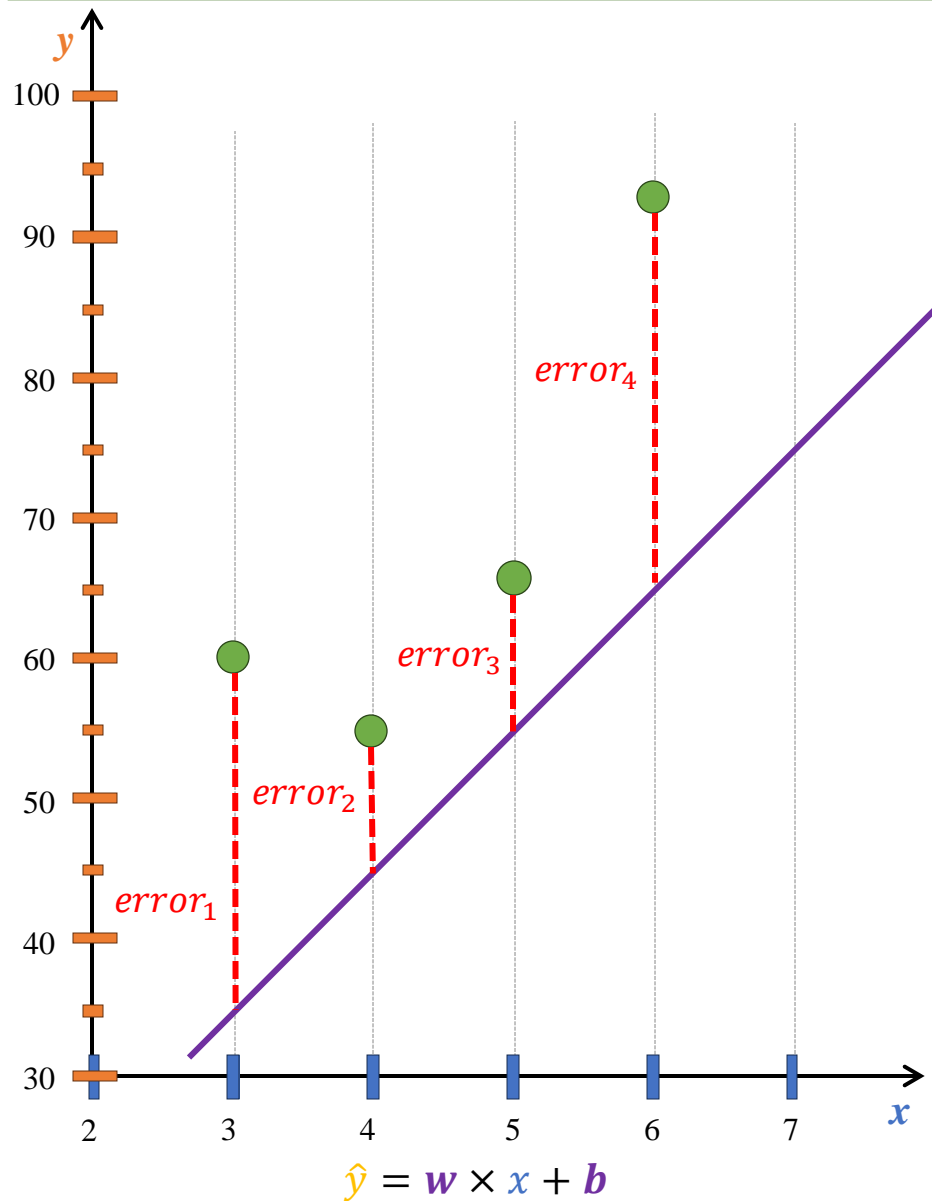
$$w = 10$$

$$b = 5$$

$$Loss = \sum_i (\hat{y}_i - y_i)^2 = \sum_i error_i$$

	Experience	Salary	Predicted	Error	Loss
1	3	60	35	625	
2	4	55	45	100	
3	5	66	55	121	
4	6	93	65	784	

Simple Linear Regression



$$\hat{y}_i = wx_i + b$$

$$error(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

$$w = 10$$

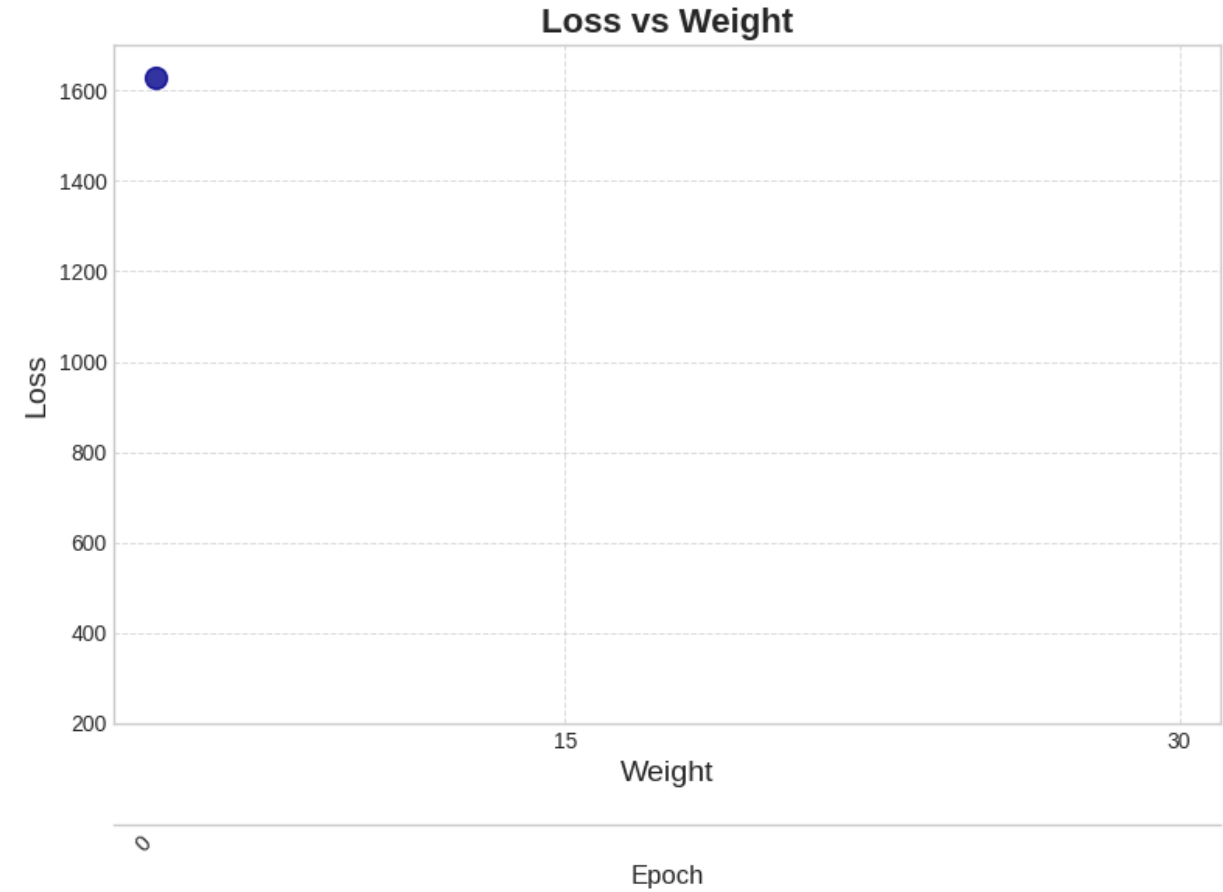

$$b = 5$$

$$Loss = \sum_i (\hat{y}_i - y_i)^2 = \sum_i error_i$$

	Experience	Salary	Predicted	Error	Loss
1	3	60	35	625	625
2	4	55	45	100	+100
3	5	66	55	121	+121
4	6	93	65	784	+784
					= 1630

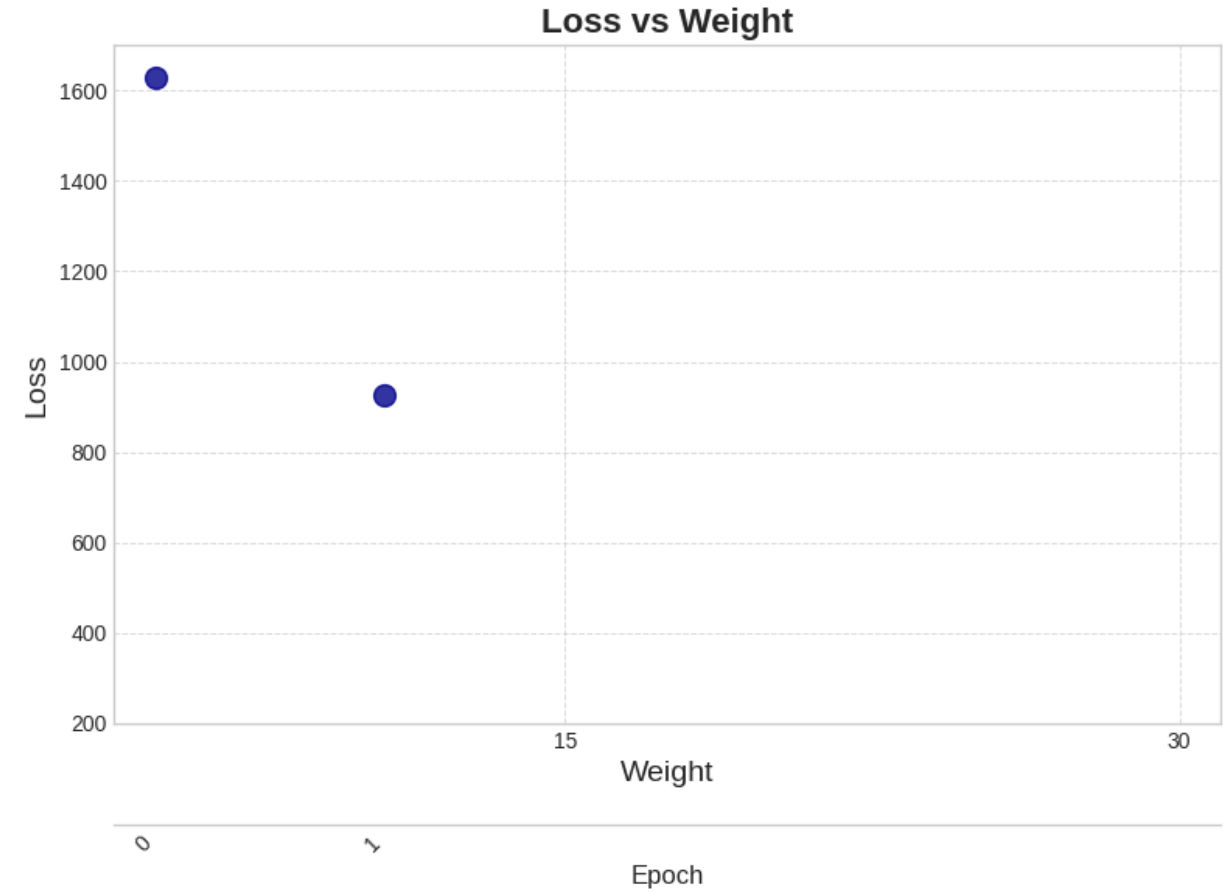

Simple Linear Regression

❖ How bias affect the loss function?





Simple Linear Regression

❖ How bias affect the loss function?




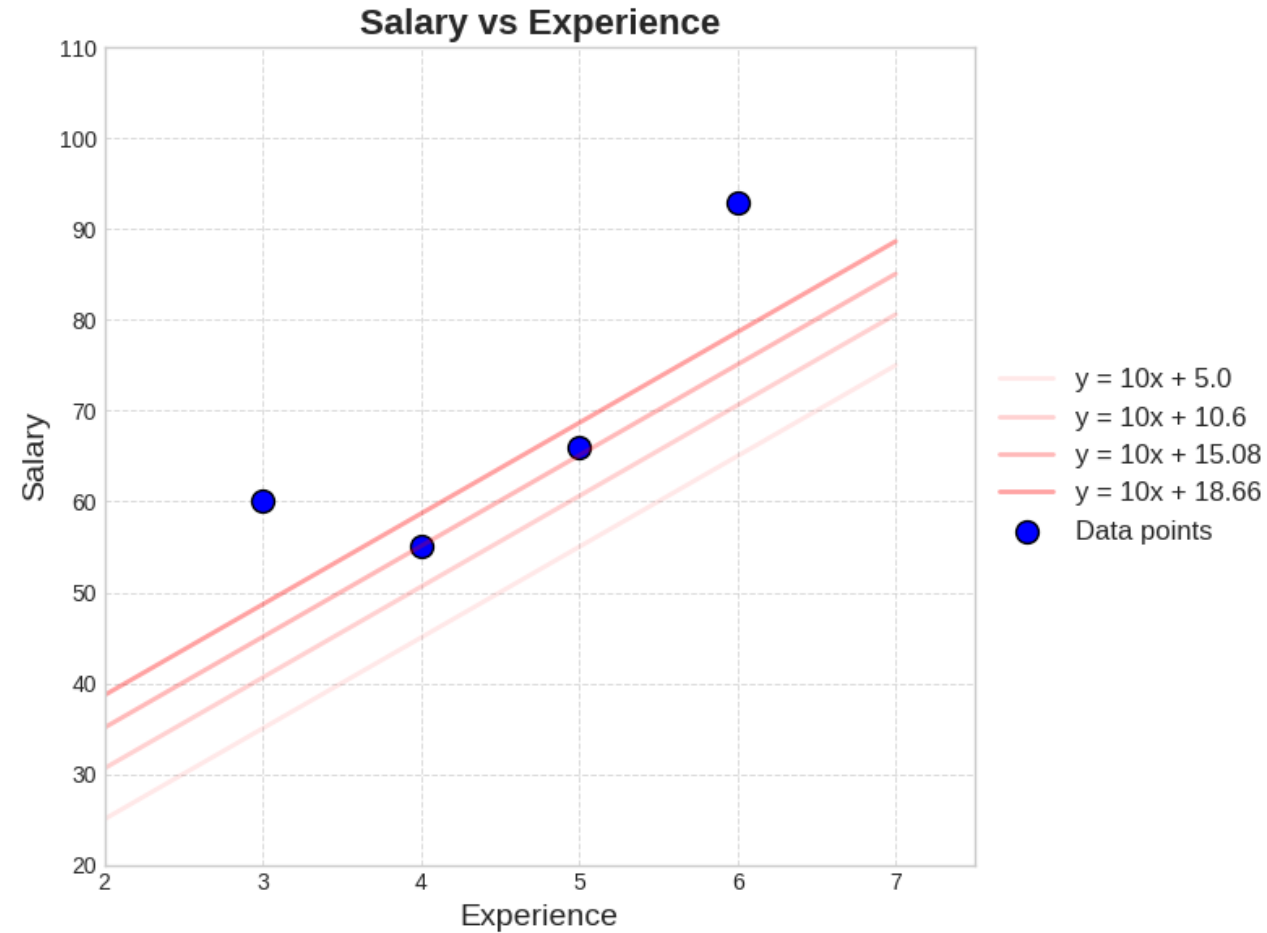
Simple Linear Regression

❖ How bias affect the loss function?





Simple Linear Regression

❖ How bias affect the loss function?




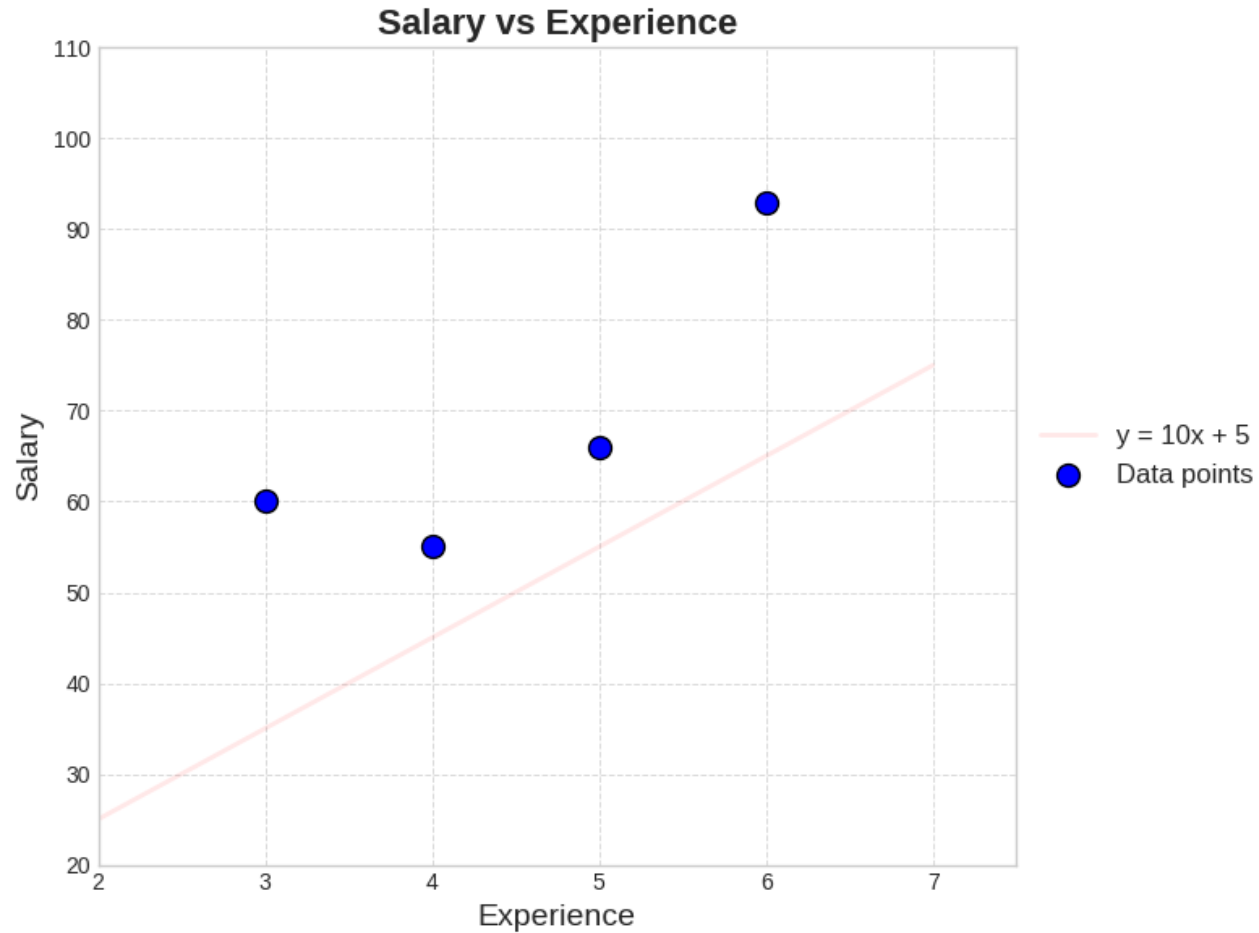
Simple Linear Regression

❖ How bias affect the loss function?





Simple Linear Regression

❖ How bias affect the loss function?





Simple Linear Regression

❖ How bias affect the loss function?




Simple Linear Regression

❖ How bias affect the loss function?





Simple Linear Regression

❖ How bias affect the loss function?




Simple Linear Regression

❖ How bias affect the loss function?



Simple Linear Regression

❖ How bias affect the loss function?




Simple Linear Regression


❖ Toward Gradient Descent




$$error_i = distance(\hat{y}_i, y_i)$$



$$\hat{y} = w_1x + b_1$$



$$\hat{y} = w_2x + b_2$$



$$\hat{y} = w_3x + b_3$$

Find a and b whose model has the smallest error, where $error = \sum_i error_i$

How?



Simple Linear Regression

❖ Toward Gradient Descent

$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

How to change w and b
so that $L(\hat{y}_i, y_i)$ reduces




Simple Linear Regression

❖ Toward Gradient Descent


$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

How to change w and b
so that $L(\hat{y}_i, y_i)$ reduces



Different b values with a fixed w value





Different w values with a fixed b value

Simple Linear Regression

❖ Toward Gradient Descent

$$\hat{y}_i = wx_i + b$$


$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$



Different b values with a fixed w value

Simple Linear Regression



❖ Derivative of a function



$$\frac{d}{dx}y(x) = \lim_{\Delta x \rightarrow 0} \frac{y(x + \Delta x) - y(x)}{\Delta x}$$



Simple Linear Regression

❖ How gradient descent works?




Simple Linear Regression

❖ How gradient descent works?



Simple Linear Regression

❖ How gradient descent works: Update gradient



$$\boxed{\frac{d}{dx} y(x)} = \lim_{\Delta x \rightarrow 0} \frac{y(x + \Delta x) - y(x)}{\Delta x}$$

$$x_{new} = x_{old} - \eta \boxed{\frac{d}{dx} y(x_{old})}$$

learning rate

Simple Linear Regression

❖ Gradient Descent

Linear equation

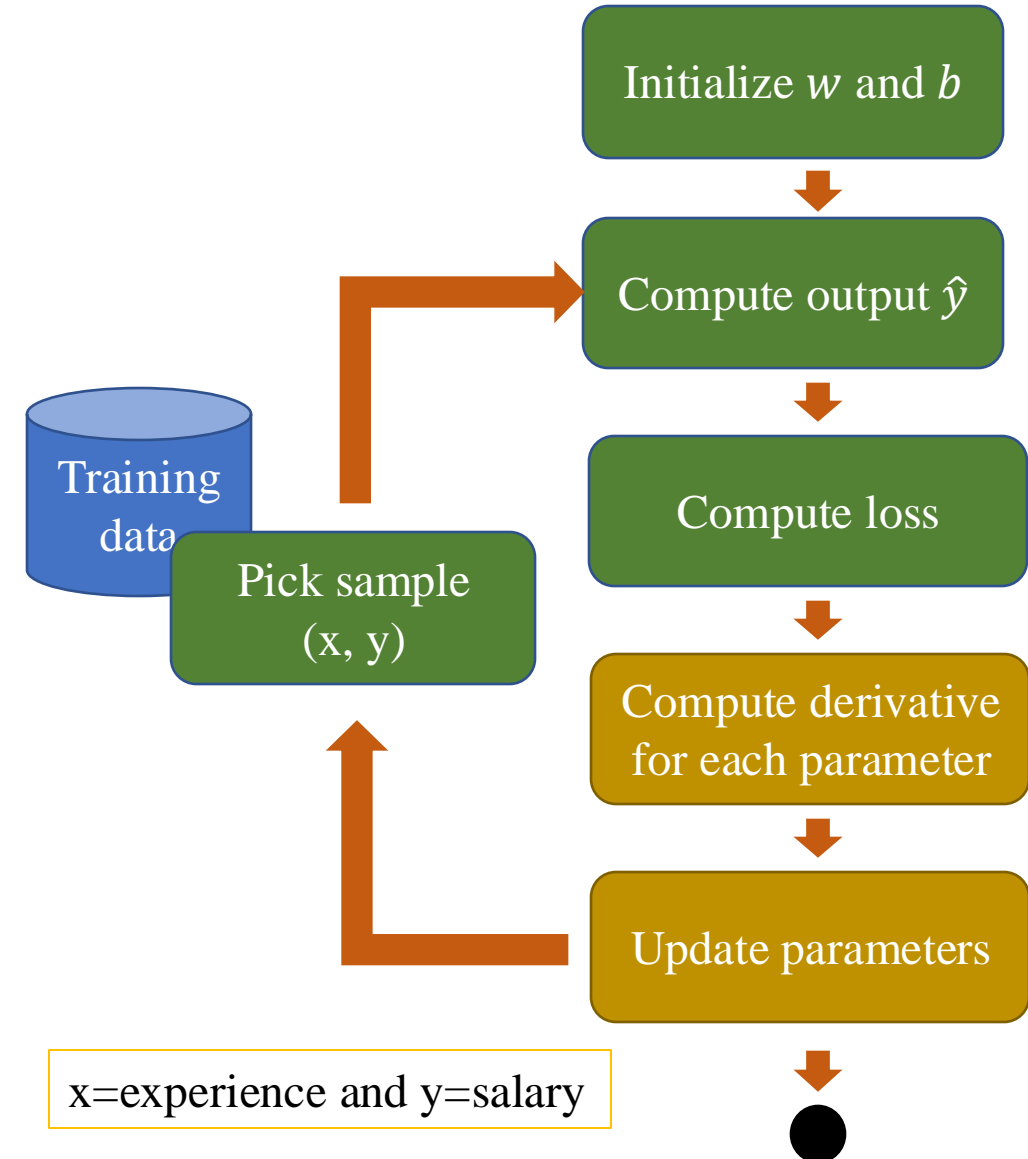
$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,
 w and b are parameters
and x is input feature

Error (loss) computation

Idea: compare predicted values \hat{y} and label values y
Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$



Simple Linear Regression

❖ Gradient Descent

Linear equation

$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,

w and b are parameters

and x is input feature

Error (loss) computation

Idea: compare predicted values \hat{y} and label values y

Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$

Find better w and b

Use gradient descent to minimize the loss function

Compute derivate for each parameter

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = 2(\hat{y} - y)$$

Update parameters

$$w = w - \eta \frac{\partial L}{\partial w} \quad b = b - \eta \frac{\partial L}{\partial b}$$

η is learning rate

Simple Linear Regression

❖ Proof: Bias (Intercept)

$$\text{Loss} = \sum_i (\hat{y} - y)^2 = \sum_i ((wx + b) - y)^2$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = \frac{\partial}{\partial \hat{y}} (\hat{y} - y)^2 \cdot \frac{\partial}{\partial b} (wx + b)$$

$$\begin{aligned} \frac{\partial L}{\partial \hat{y}} &= \frac{\partial}{\partial \hat{y}} (\hat{y} - y)^2 = 2 \cdot (\hat{y} - y) \cdot \frac{\partial (\hat{y}_i - y_i)}{\partial \hat{y}_i} \\ &= 2 \cdot (\hat{y} - y) \cdot \frac{\partial}{\partial \hat{y}_i} (\hat{y}_i - y_i) \\ &= 2 \cdot (\hat{y} - y) \cdot 1 \\ &= 2(\hat{y} - y) \end{aligned}$$

$$\begin{aligned} \frac{\partial \hat{y}}{\partial b} &= \frac{\partial}{\partial b} (wx + b) \\ &= 1 \end{aligned}$$

$$\Rightarrow \frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = 2(\hat{y} - y)$$

Simple Linear Regression

❖ Update Bias (Intercept)

$$x_{new} = x_{old} - \eta \frac{d}{dx} y(x_{old})$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = 2(\hat{y} - y)$$

$$b_{new} = b_{old} - \eta \frac{\partial L}{\partial b}$$

$$b_{new} = b_{old} - \eta 2(\hat{y} - y)$$

Simple Linear Regression

❖ Proof: Slope

$$\text{Loss} = \sum_i (\hat{y} - y)^2 = \sum_i ((wx + b) - y)^2$$

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} = \frac{\partial}{\partial \hat{y}} (\hat{y} - y)^2 \cdot \frac{\partial}{\partial w} (wx + b)$$

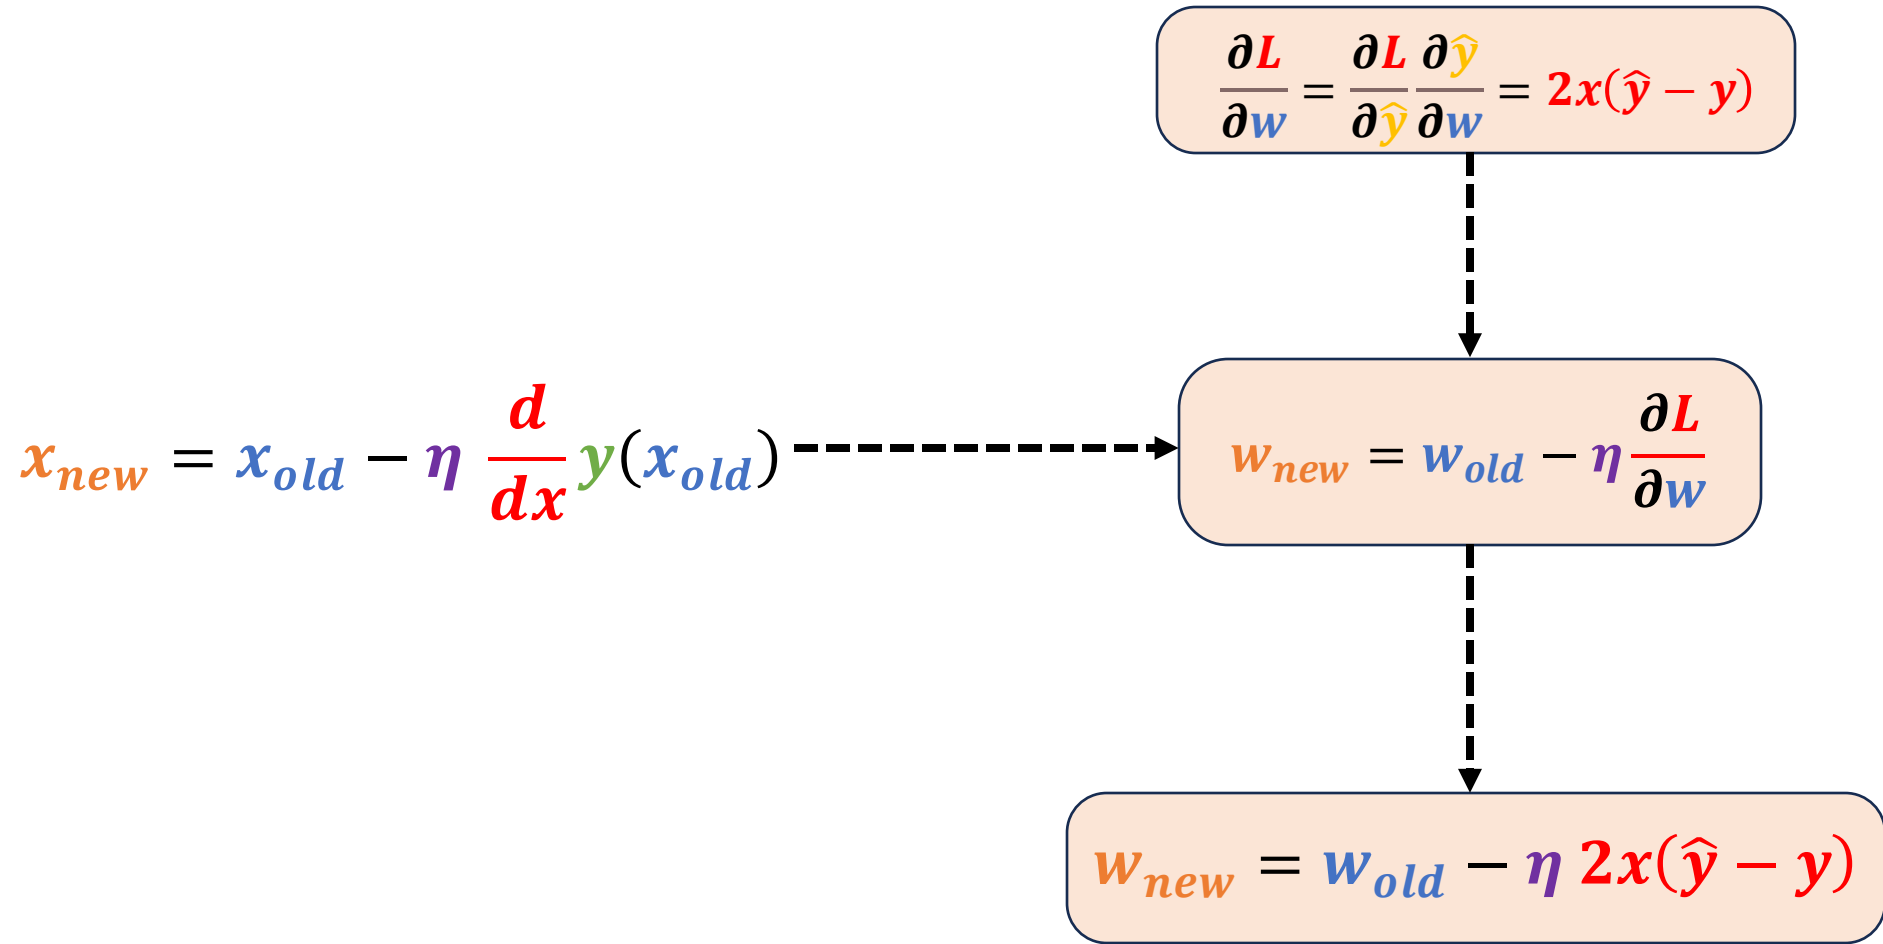
$$\begin{aligned} \frac{\partial L}{\partial \hat{y}} &= \frac{\partial}{\partial \hat{y}} (\hat{y} - y)^2 = 2 \cdot (\hat{y} - y) \cdot \frac{\partial (\hat{y}_i - y_i)}{\partial \hat{y}_i} \\ &= 2 \cdot (\hat{y} - y) \cdot \frac{\partial}{\partial \hat{y}_i} (\hat{y}_i - y_i) \\ &= 2 \cdot (\hat{y} - y) \cdot 1 \\ &= 2(\hat{y} - y) \end{aligned}$$

$$\begin{aligned} \frac{\partial \hat{y}}{\partial w} &= \frac{\partial}{\partial w} (wx + b) \\ &= x \end{aligned}$$

$$\Rightarrow \frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} = 2x(\hat{y} - y)$$


Simple Linear Regression

❖ Update Slope



Simple Linear Regression

❖ Example



$$\hat{y} = (wx + b)$$

$$\text{Loss} = (\hat{y} - y)^2$$

$$w_{\text{new}} = w_{\text{old}} - \eta \frac{\partial L}{\partial w}$$

$$b_{\text{new}} = b_{\text{old}} - \eta \frac{\partial L}{\partial b}$$

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

Simple Linear Regression

❖ Example: Forward 1

Initialize

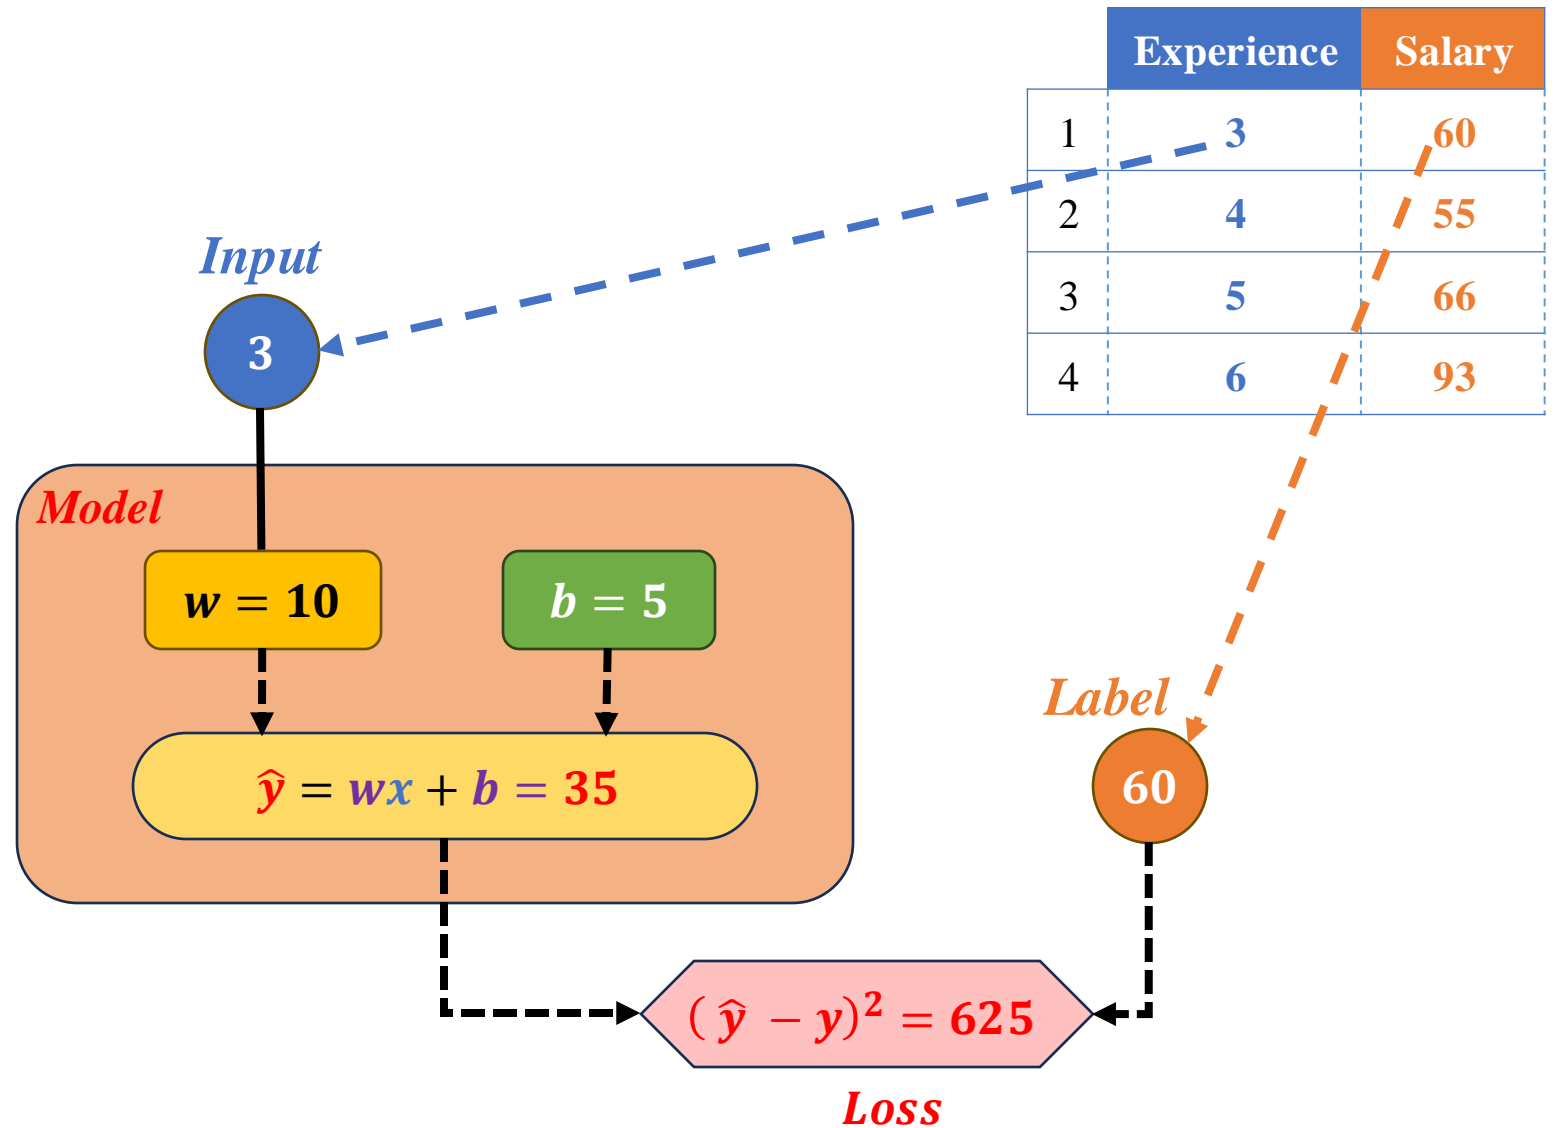
$$w = 10$$

$$b = 5$$

$$\eta = 0.01$$

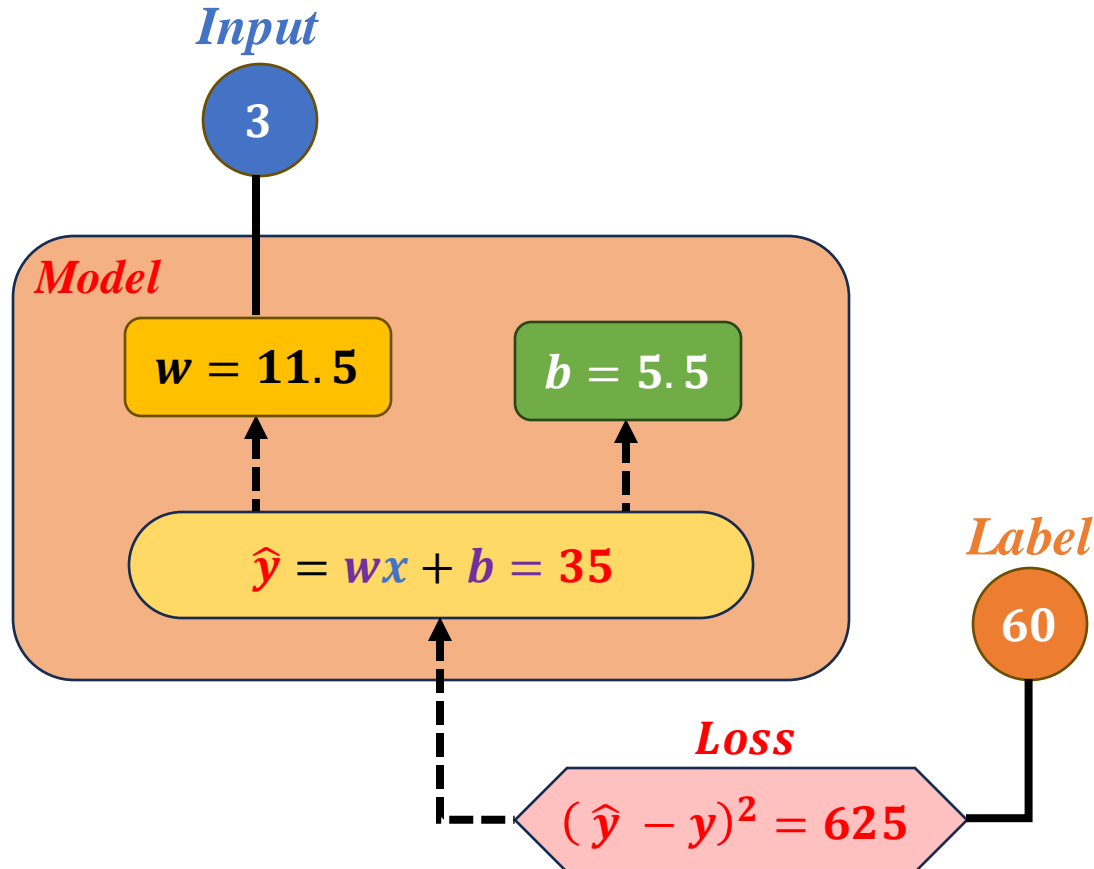
$$\hat{y} = (wx + b)$$

$$\text{Loss} = (\hat{y} - y)^2$$



Simple Linear Regression

❖ Example: Backpropagation 1



$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

$$\frac{\partial L}{\partial w} = 2 \times 3(35 - 60) = -150$$

$$\frac{\partial L}{\partial b} = 2(35 - 60) = -50$$

$$w_{new} = w_{old} - \eta \frac{\partial L}{\partial w}$$

$$11.5 = 10 - 0.01 \times (-150)$$

$$b_{new} = b_{old} - \eta \frac{\partial L}{\partial b}$$

$$5.5 = 5 - 0.01 \times (-50)$$

Simple Linear Regression

❖ Example: Forward 2

Initialize

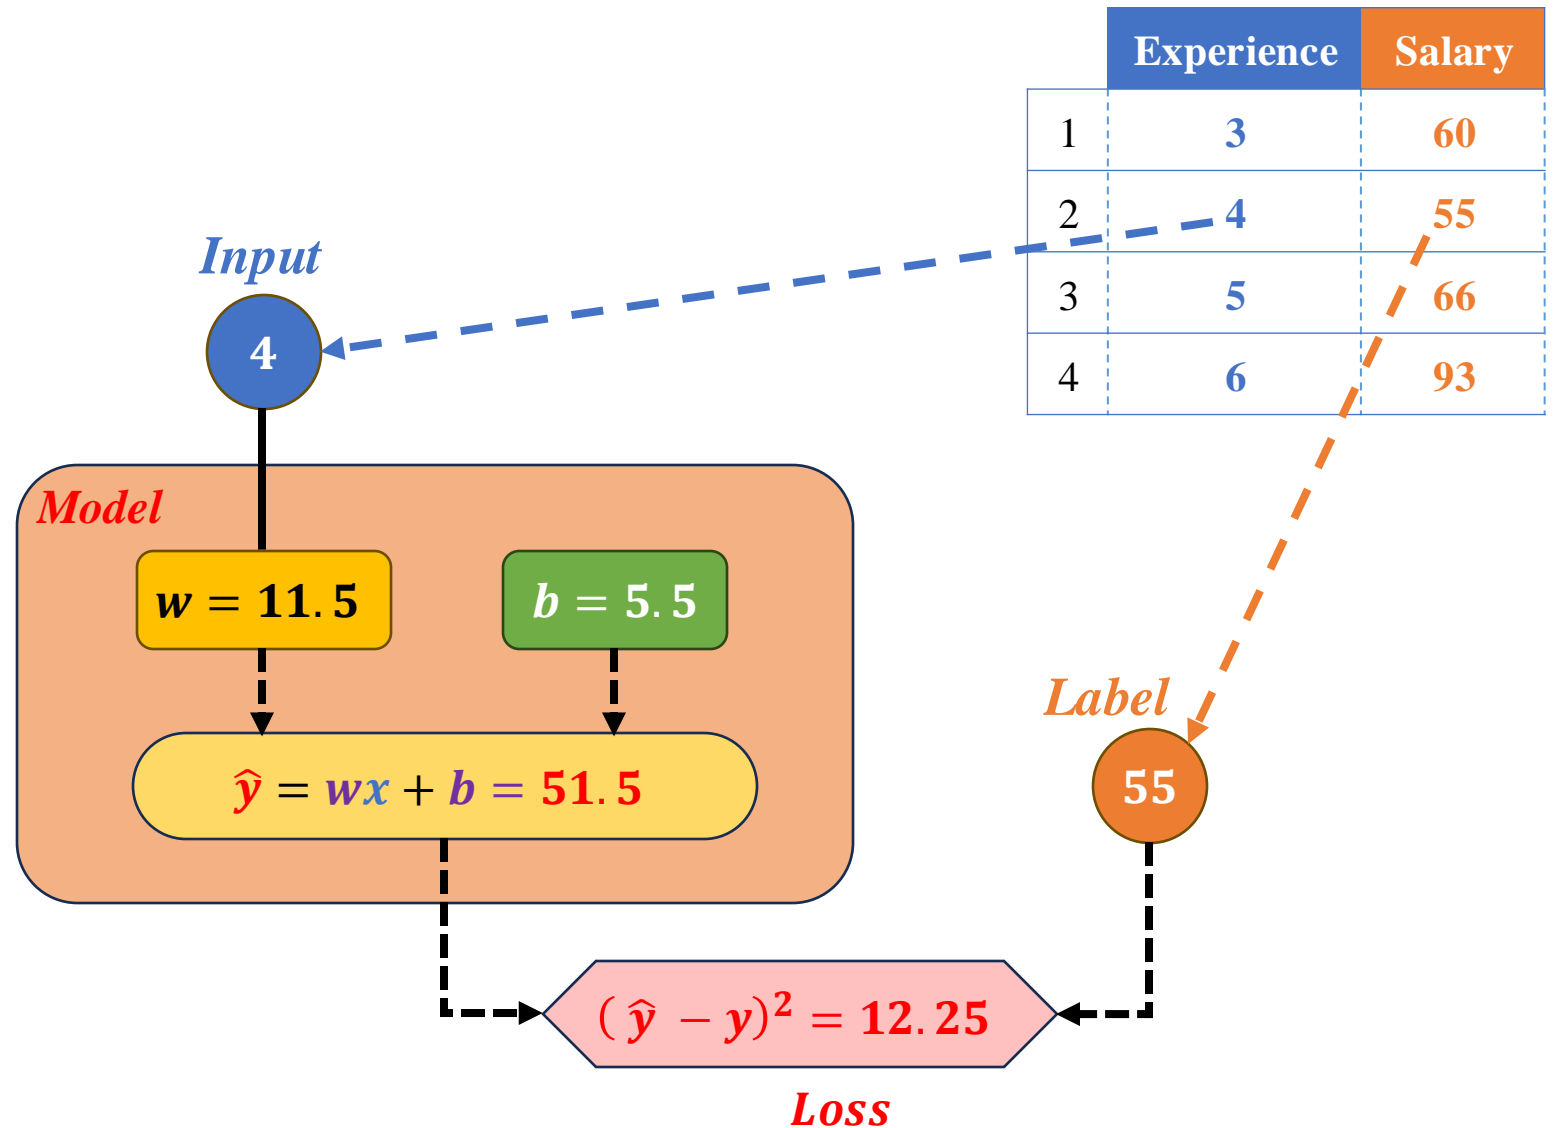
$$w = 10$$

$$b = 5$$

$$\eta = 0.1$$

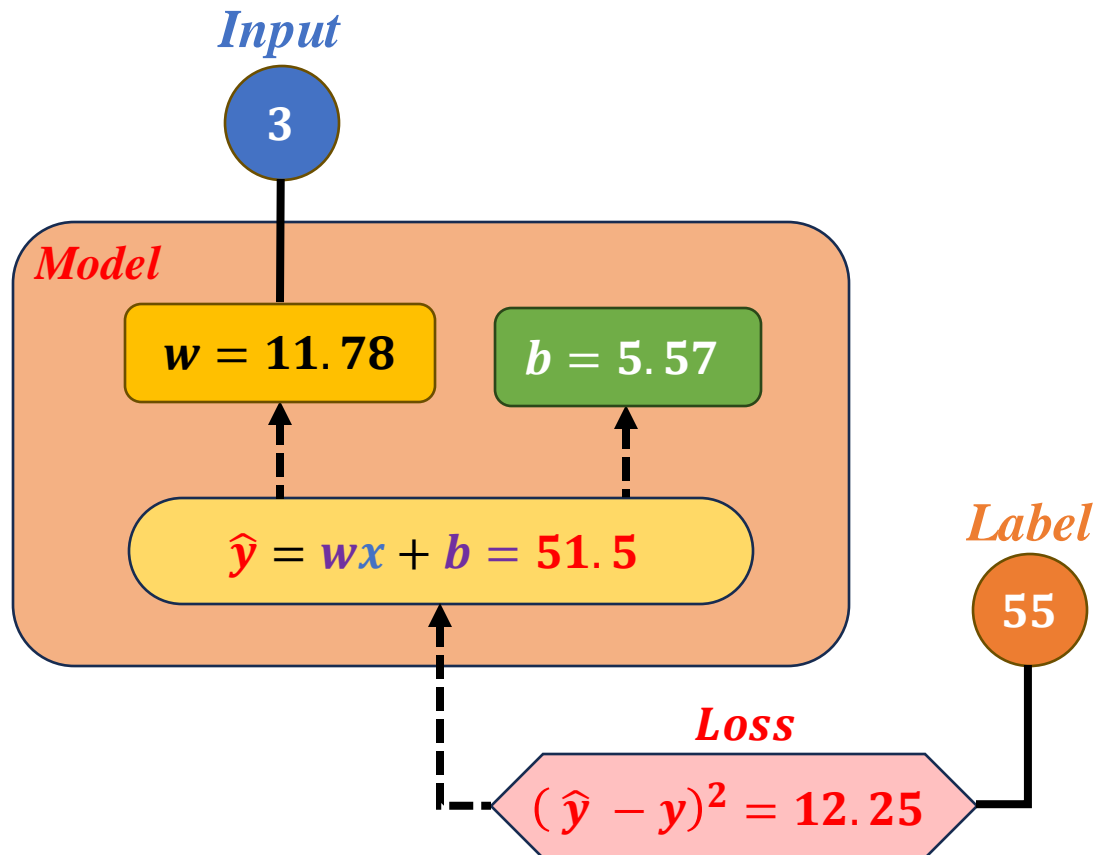
$$\hat{y} = (wx + b)$$

$$\text{Loss} = (\hat{y} - y)^2$$



Simple Linear Regression

❖ Example: Backpropagation 2



$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

$$\frac{\partial L}{\partial w} = 2 \times 3(51.5 - 55) = -28$$

$$\frac{\partial L}{\partial b} = 2(51.5 - 55) = -7$$

$$w_{new} = w_{old} - \eta \frac{\partial L}{\partial w}$$

$$b_{new} = b_{old} - \eta \frac{\partial L}{\partial b}$$


$$11.78 = 11.5 - 0.01 \times (-28)$$

$$5.57 = 5.5 - 0.01 \times (-7)$$

Simple Linear Regression

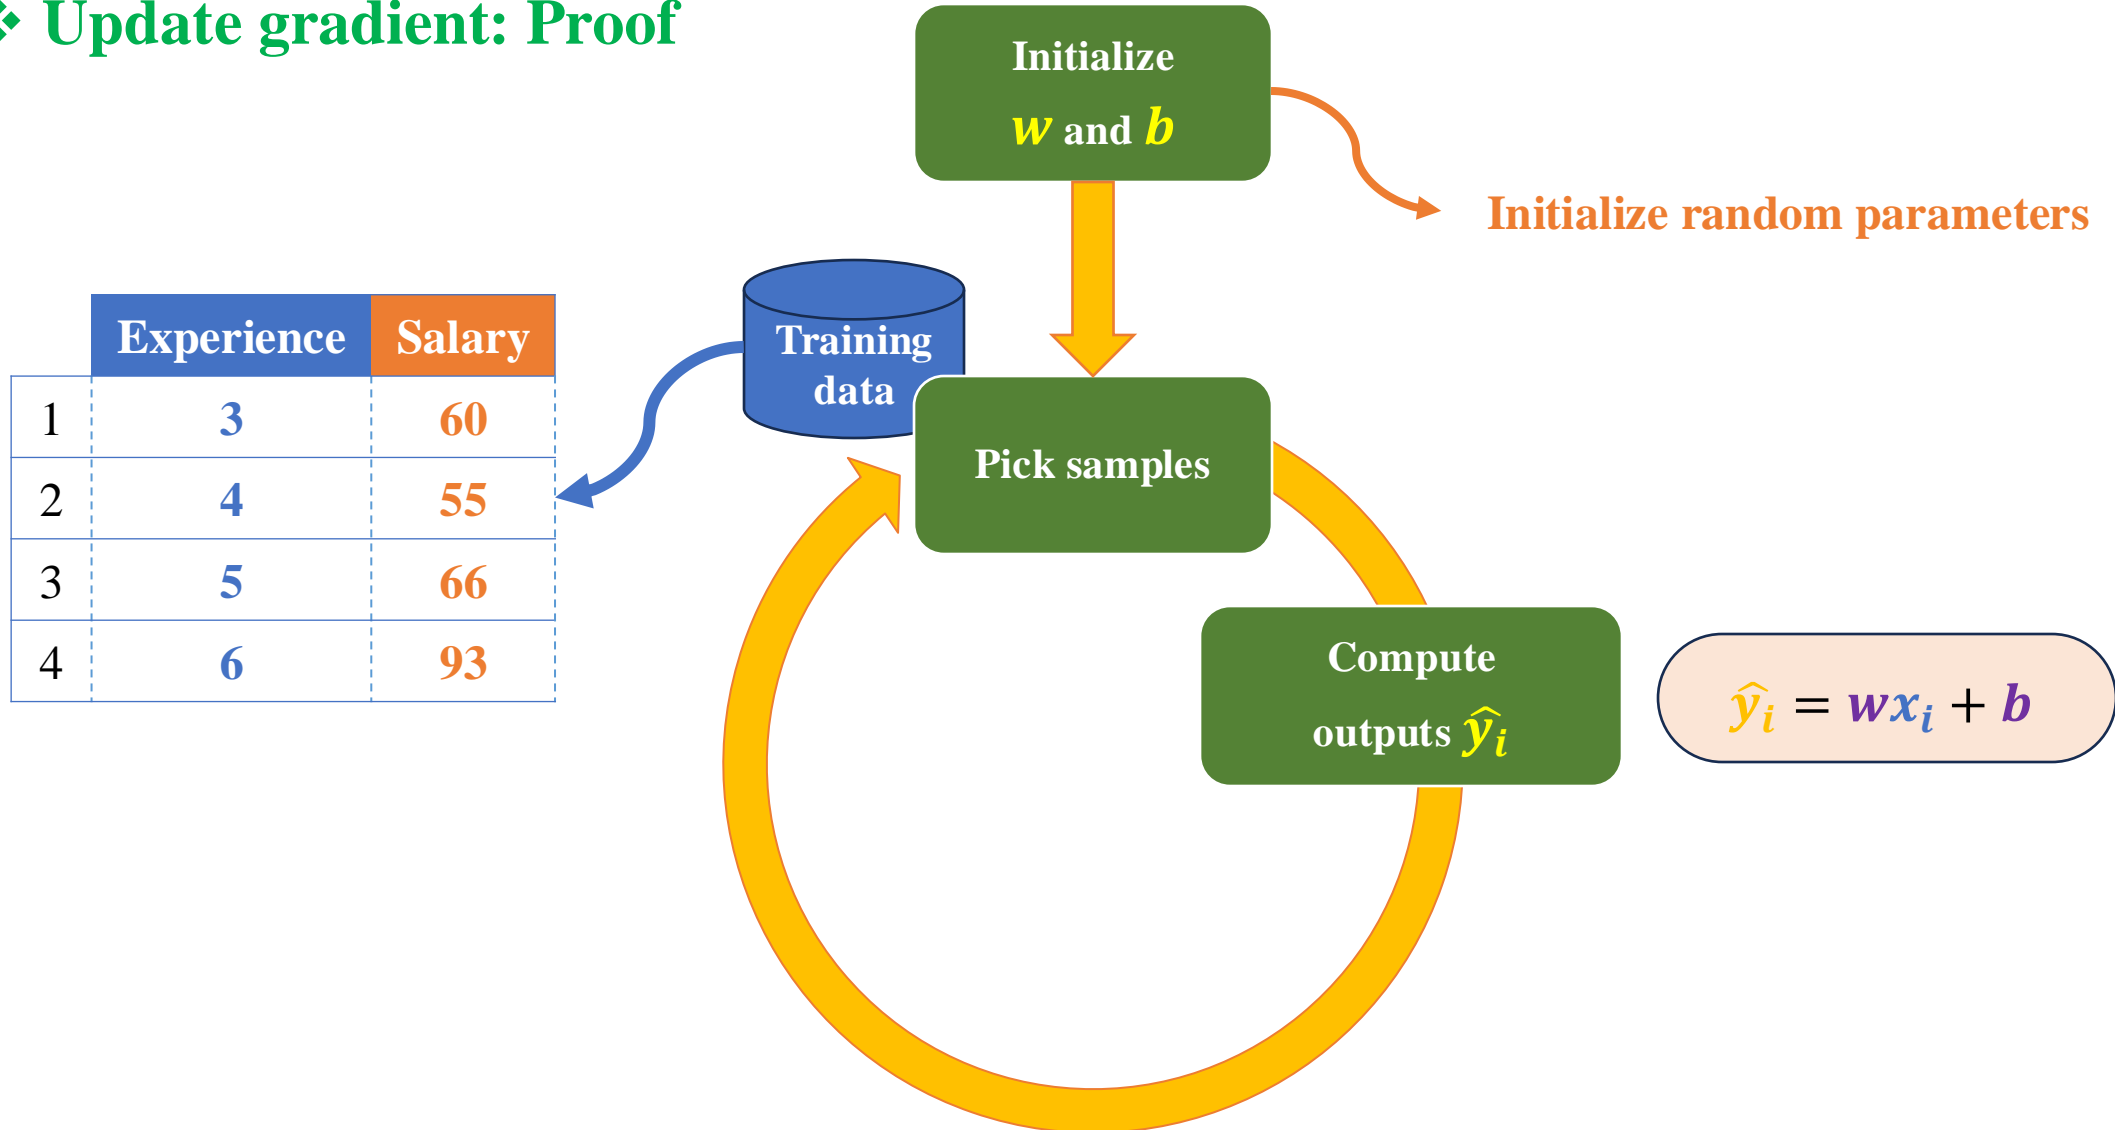
❖ Update gradient: Proof

	Experience	Salary
1	3	60
2	4	55
3	5	66
4	6	93



Simple Linear Regression

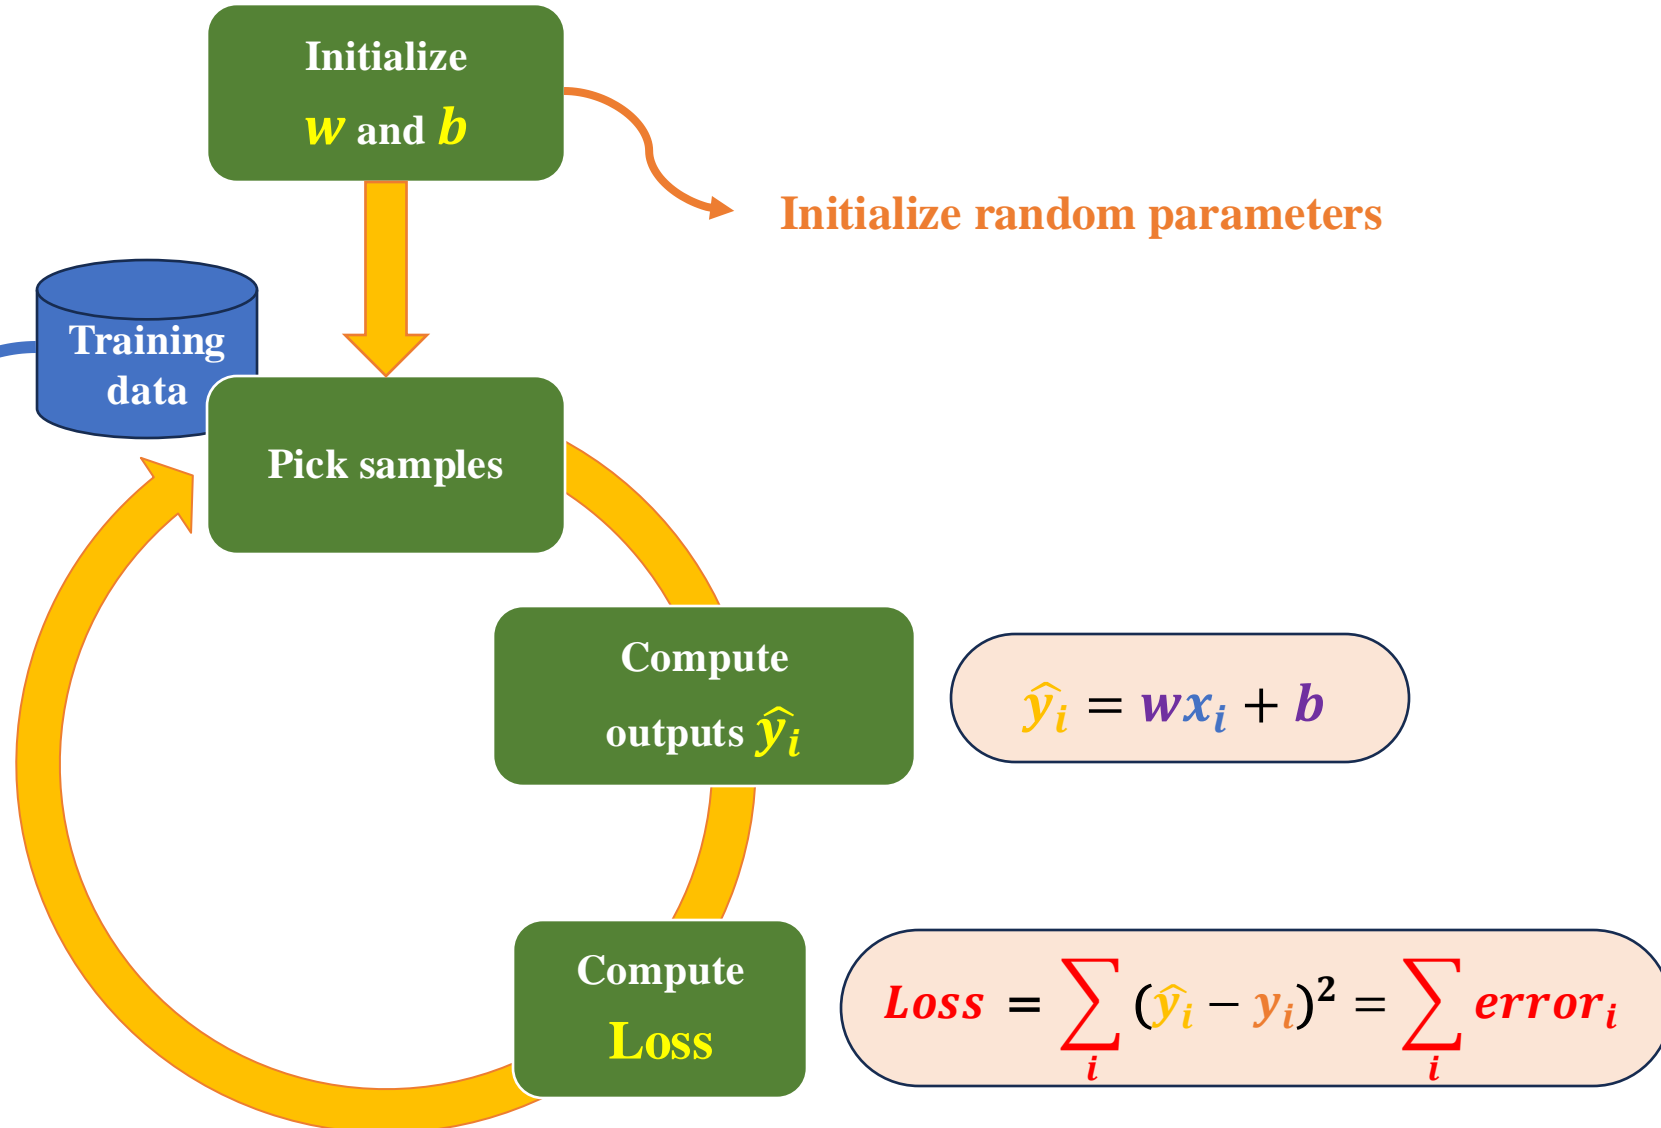
❖ Update gradient: Proof



Simple Linear Regression

❖ Update gradient: Proof

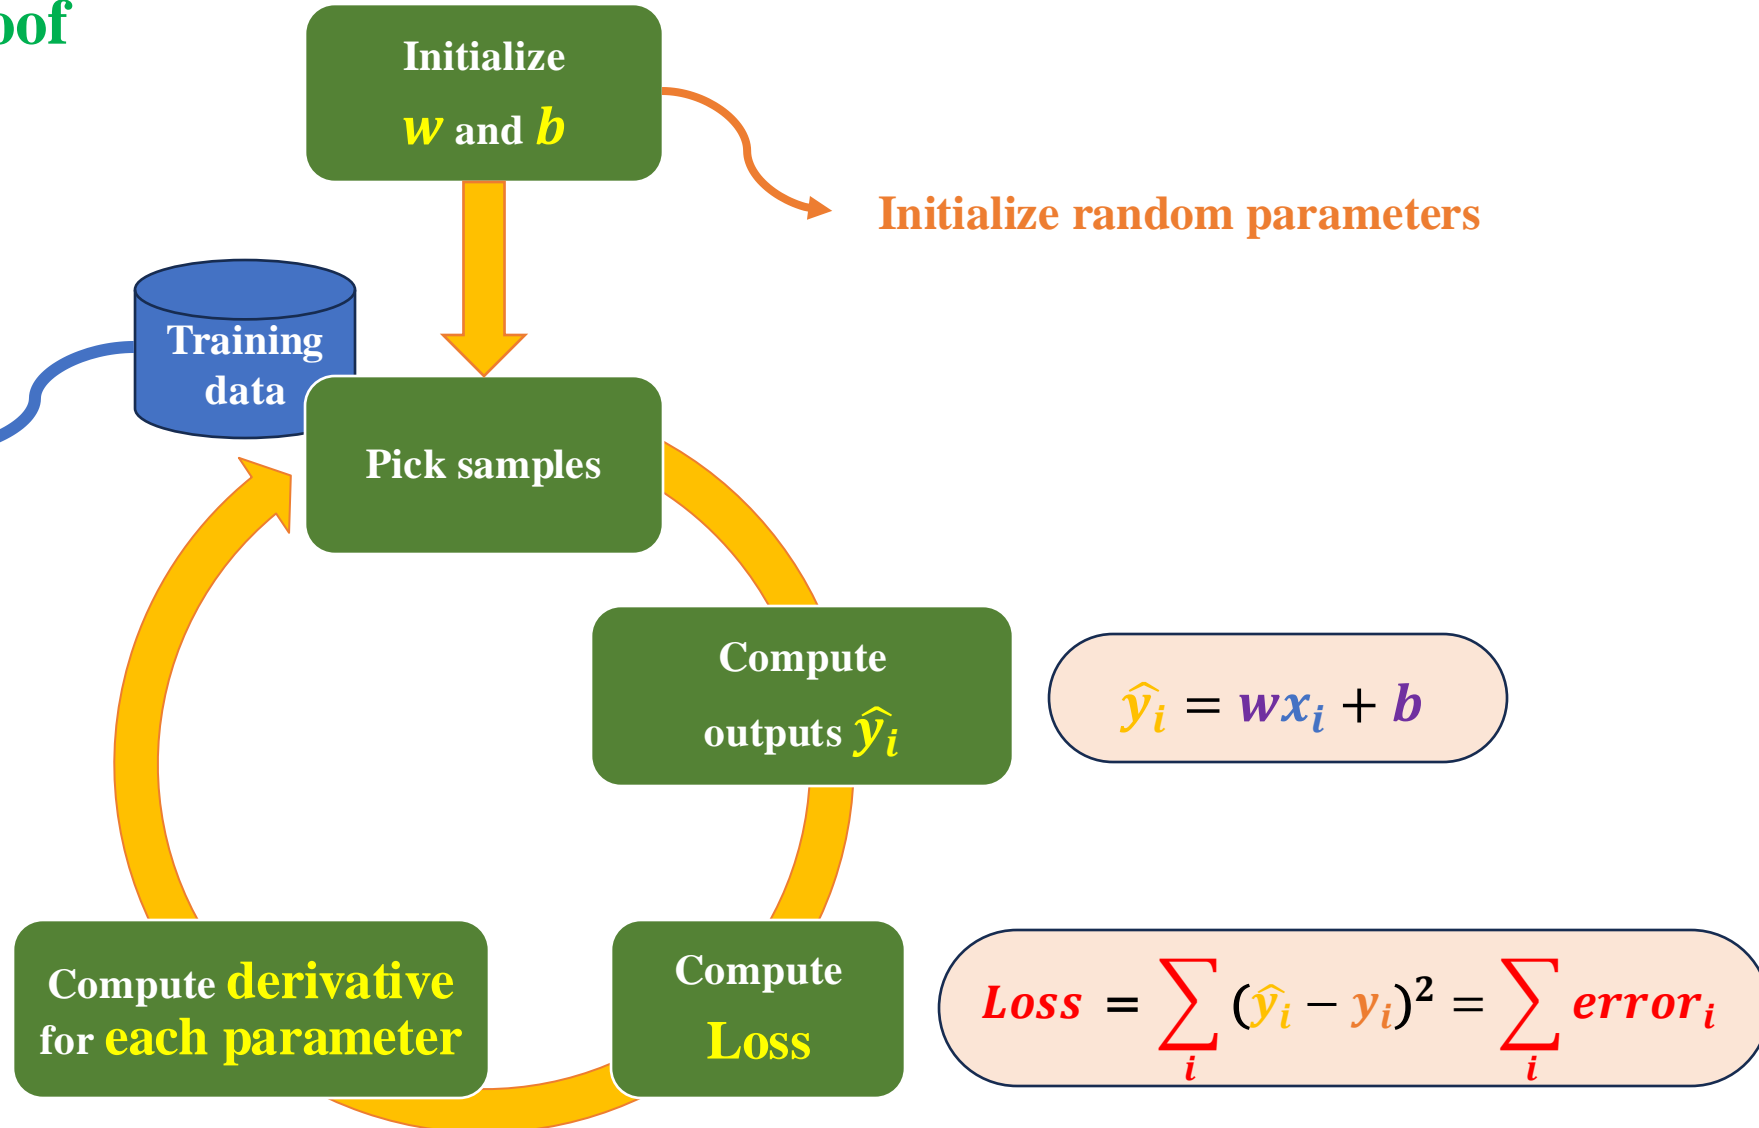
	Experience	Salary
1	3	60
2	4	55
3	5	66
4	6	93



Simple Linear Regression

❖ Update gradient: Proof

	Experience	Salary
1	3	60
2	4	55
3	5	66
4	6	93



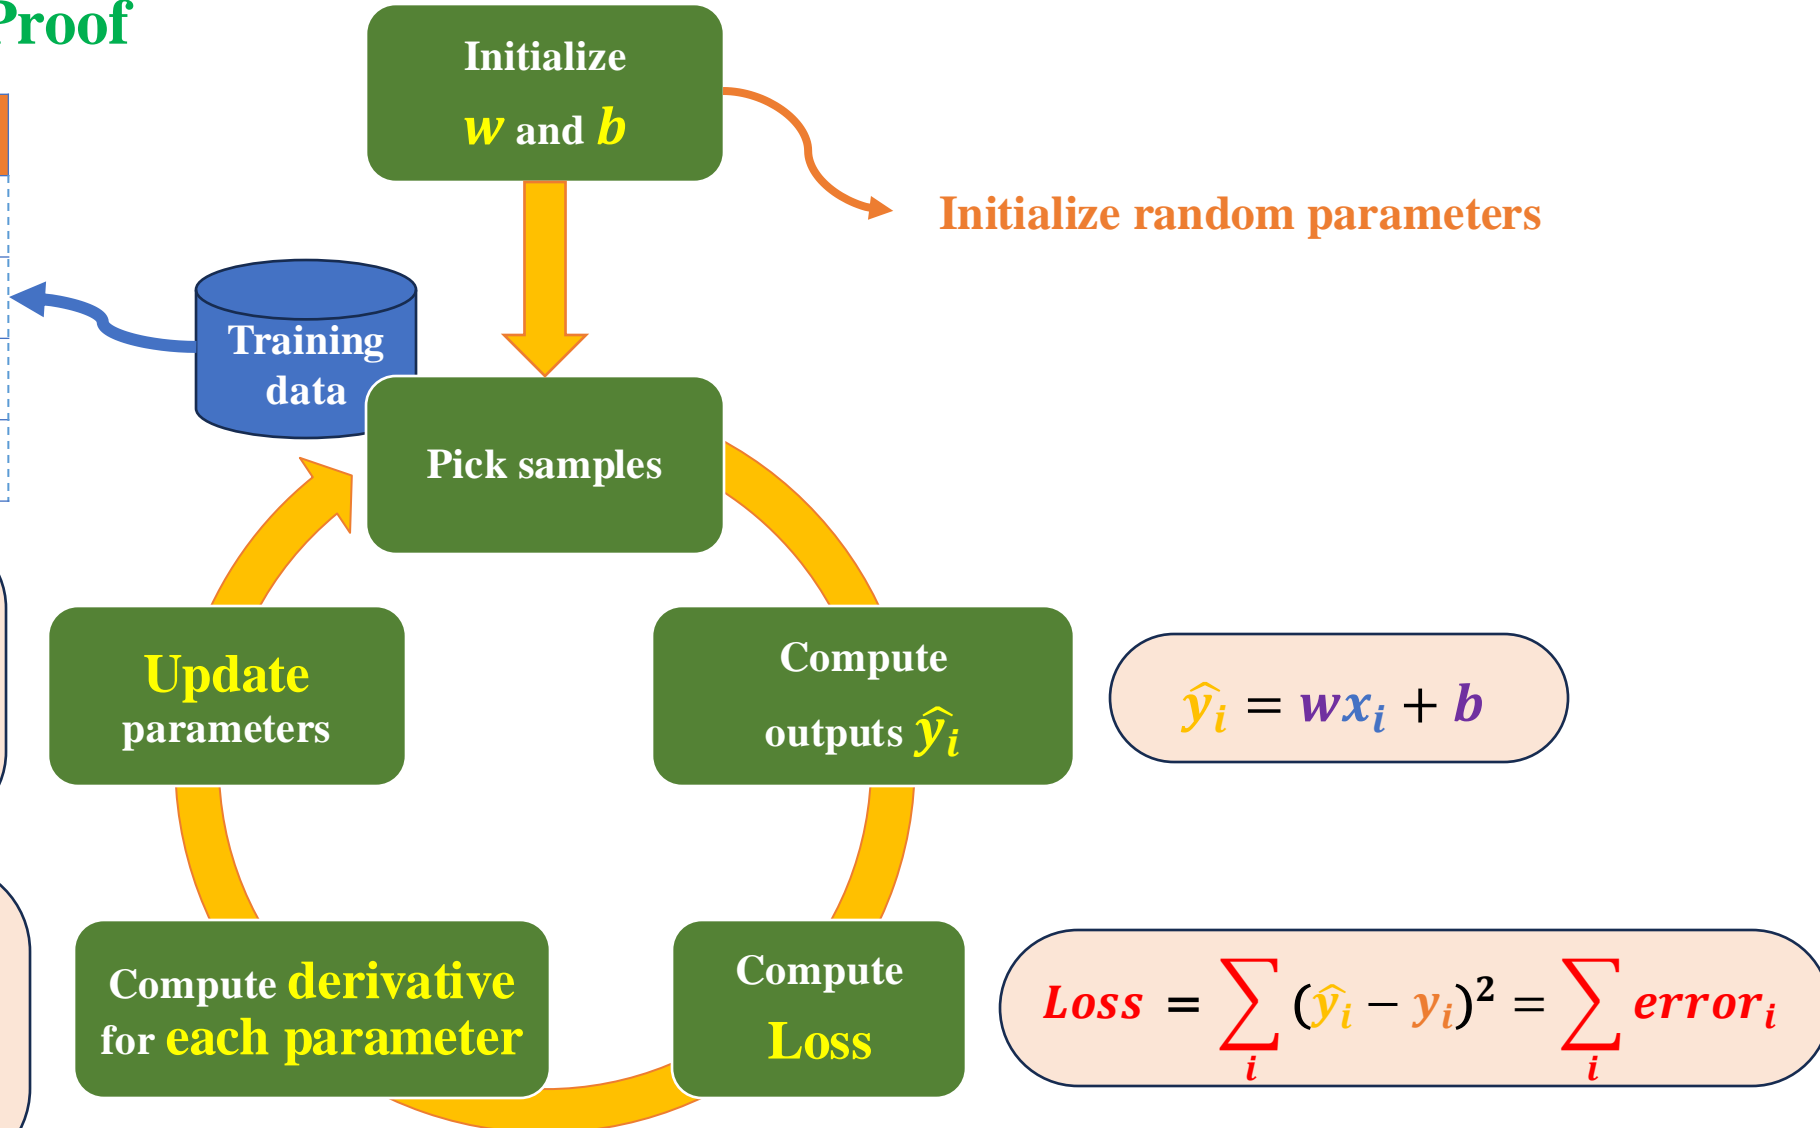
Simple Linear Regression

❖ Update gradient: Proof

	Experience	Salary
1	3	60
2	4	55
3	5	66
4	6	93

$$w_{new} = w_{old} - \eta \frac{\partial L}{\partial w}$$
$$b_{new} = b_{old} - \eta \frac{\partial L}{\partial b}$$

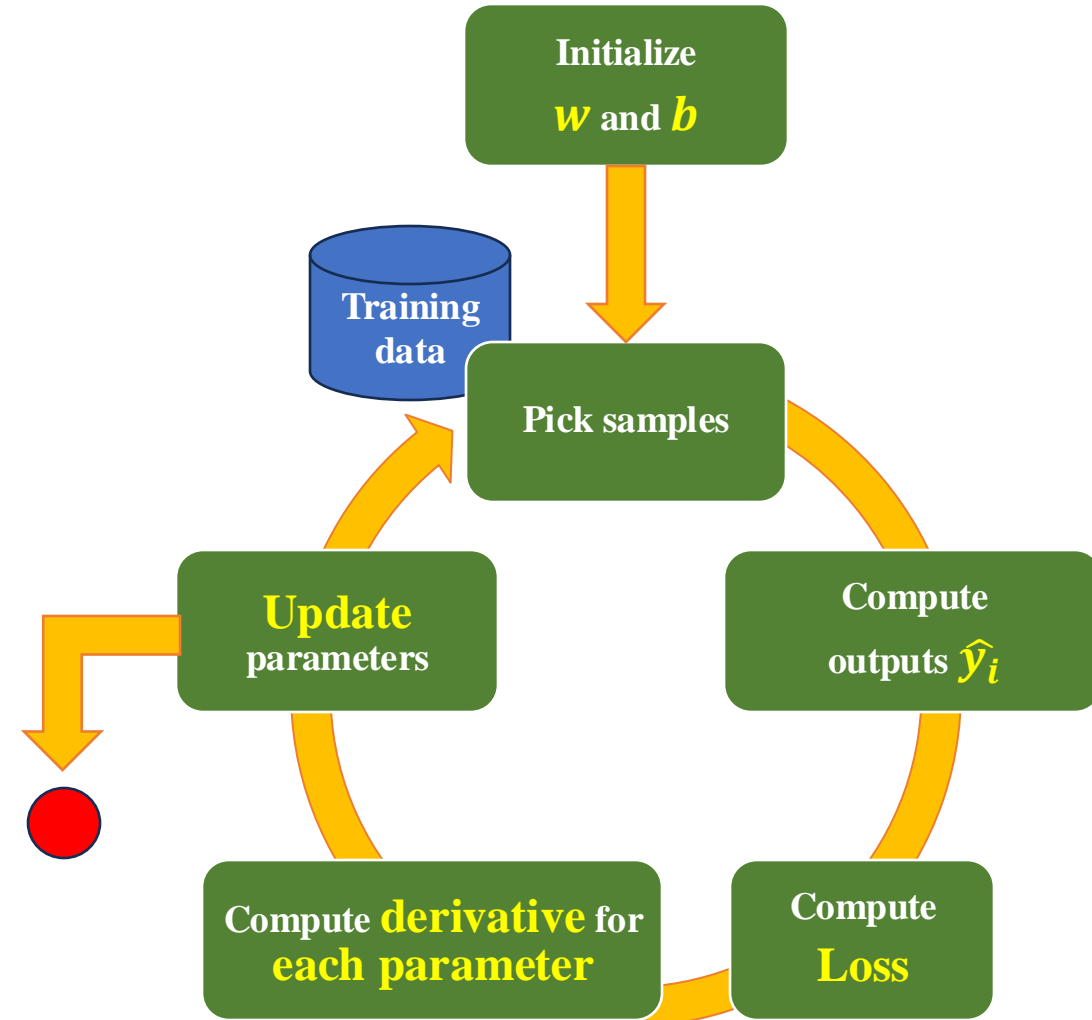

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = 2(\hat{y} - y)$$
$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} = 2x(\hat{y} - y)$$



Simple Linear Regression

❖ Update gradient: Proof

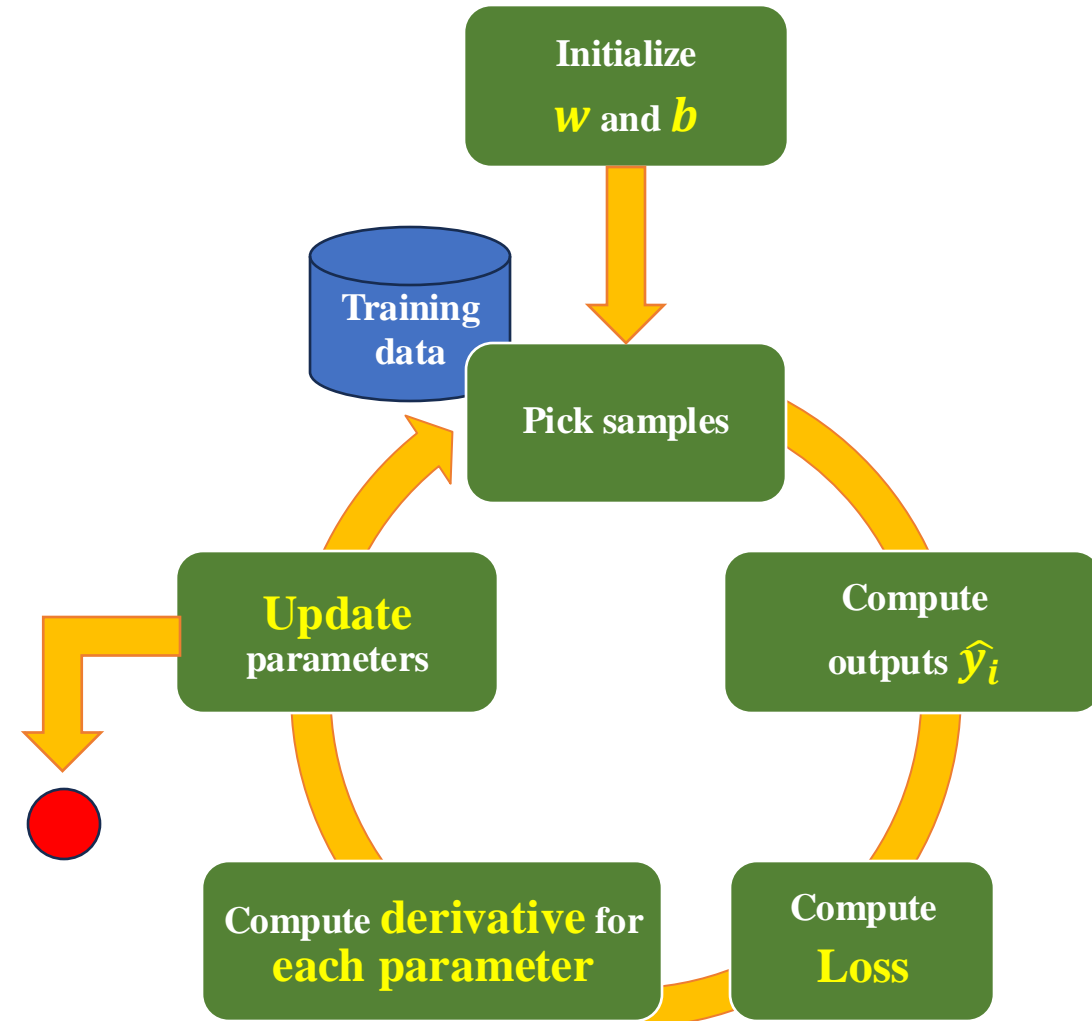

Salary vs Experience



Simple Linear Regression

❖ Update gradient: Proof

Loss vs Weight and Bias

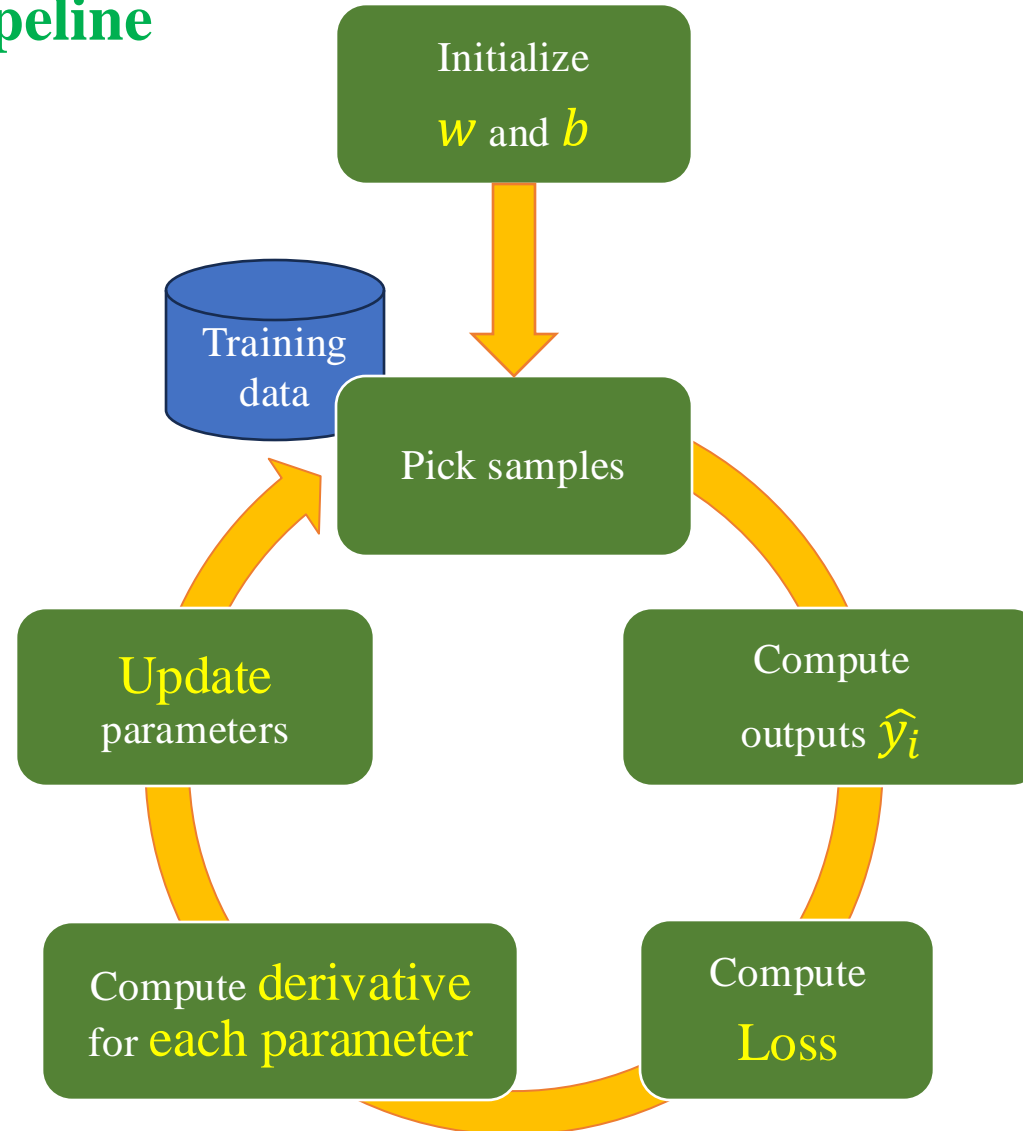


QUIZ

Code Implementation

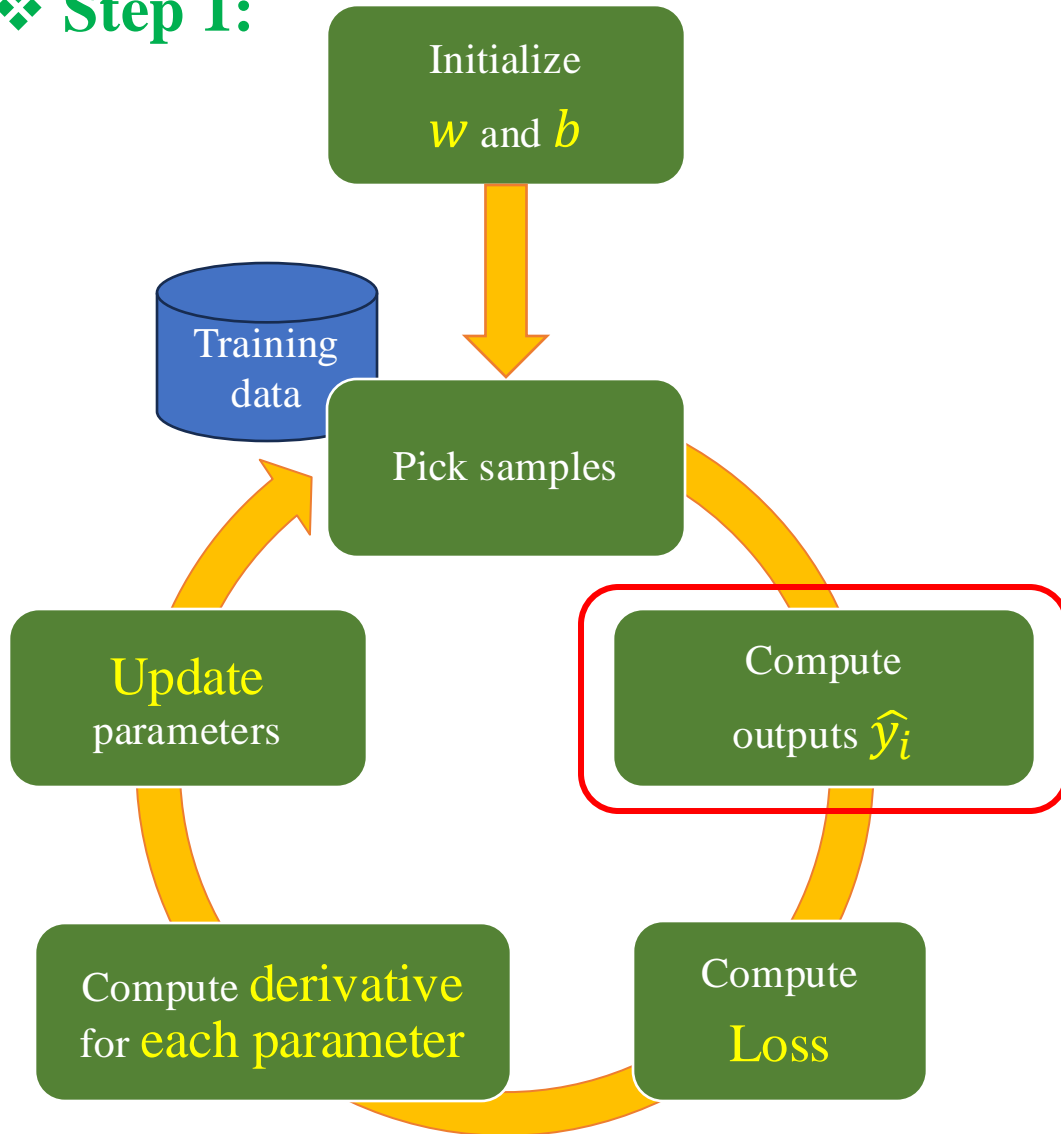
Code Implementation

❖ Gradient Descent Pipeline



Code Implementation

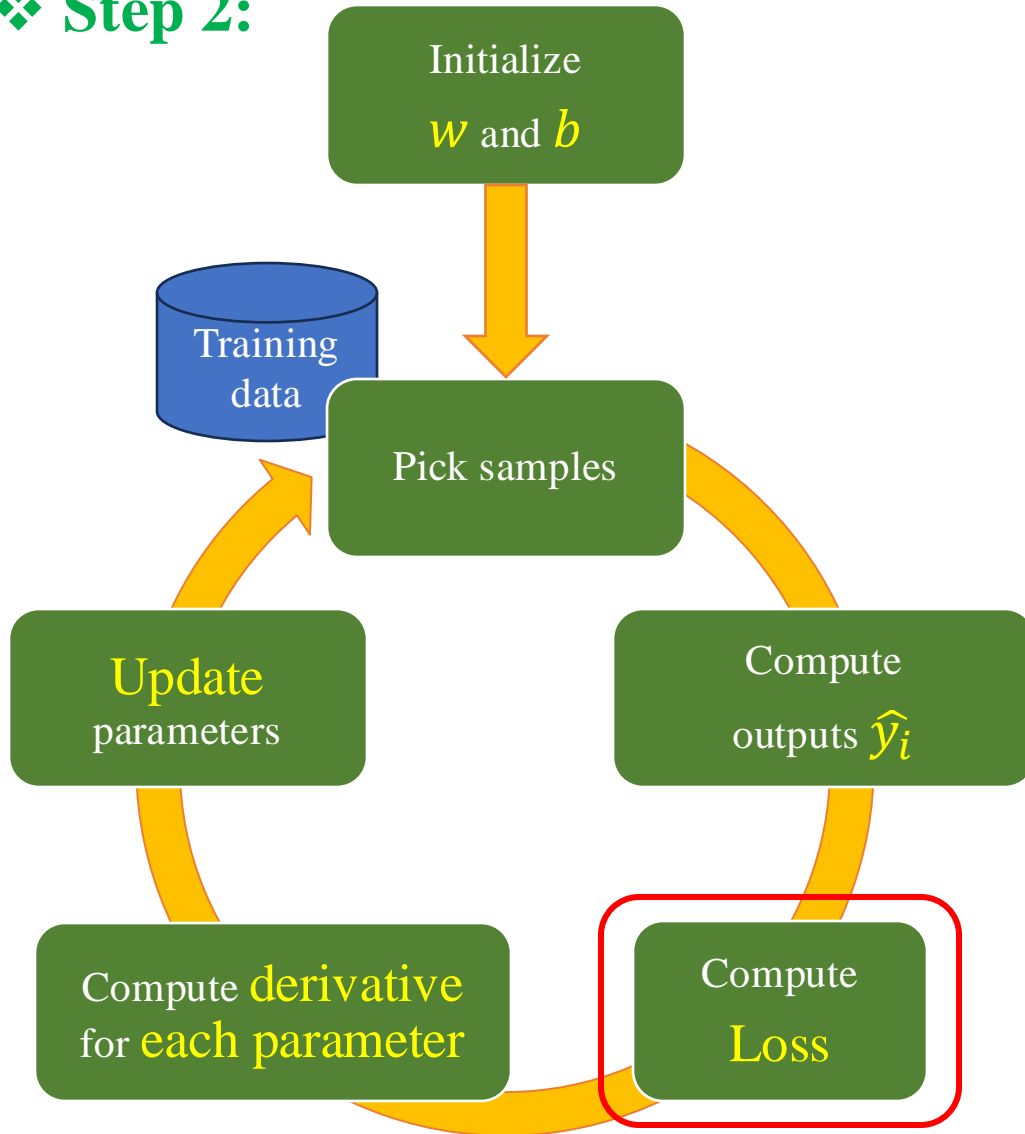
❖ Step 1:



```
1 def predict(x, w, b):  
2     return x * w + b  
3  
4 def compute_loss(y_hat, y):  
5     return (y_hat - y) ** 2  
6  
7 def compute_gradient(y_hat, y, x):  
8     dw = 2 * x * (y_hat - y)  
9     db = 2 * (y_hat - y)  
10  
11     return dw, db  
12  
13 def update_weight(w, b, dw, db, lr):  
14     new_w = w - lr * dw  
15     new_b = b - lr * db  
16  
17     return new_w, new_b
```

Code Implementation

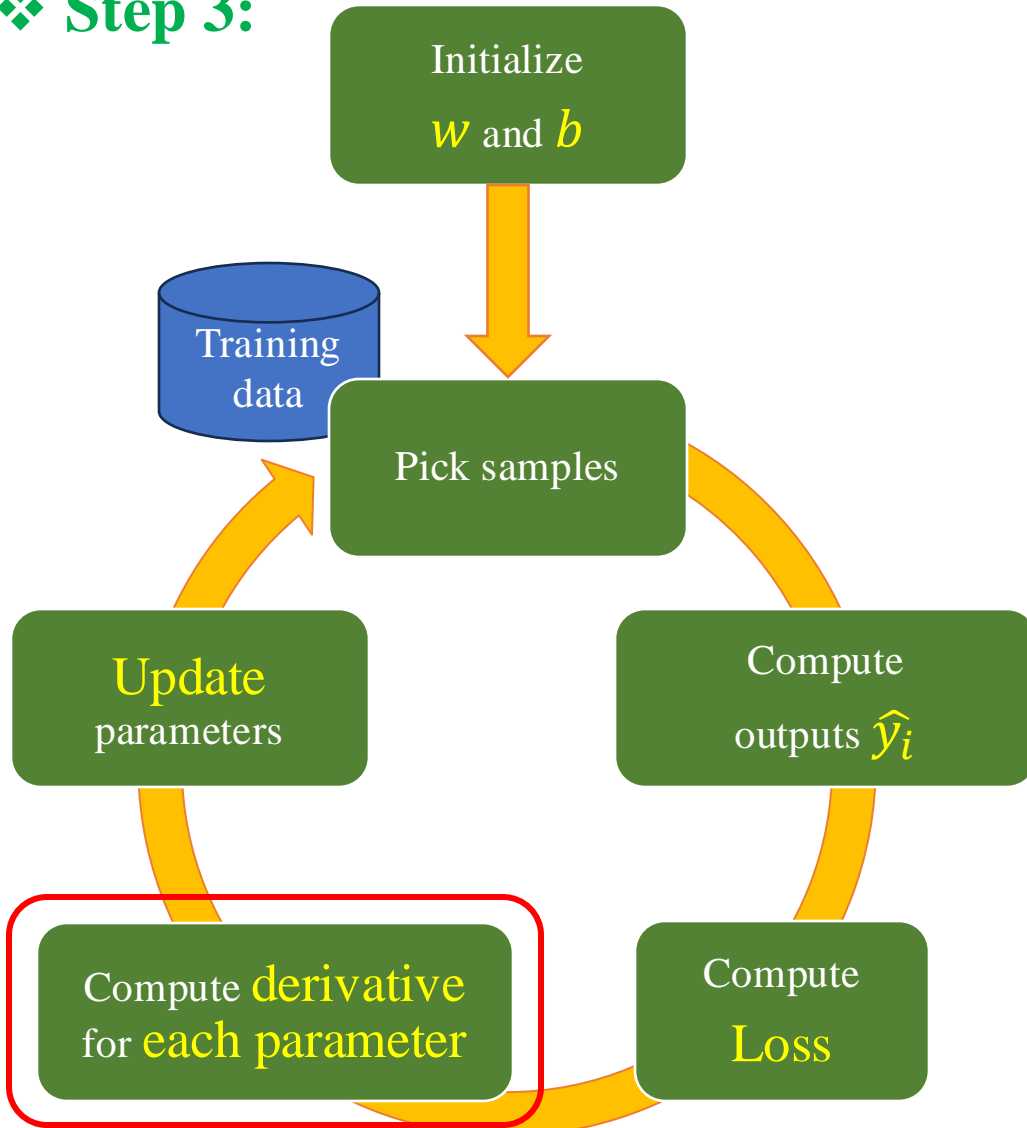
❖ Step 2:



```
1 def predict(x, w, b):
2     return x * w + b
3
4 def compute_loss(y_hat, y):
5     return (y_hat - y) ** 2
6
7 def compute_gradient(y_hat, y, x):
8     dw = 2 * x * (y_hat - y)
9     db = 2 * (y_hat - y)
10
11     return dw, db
12
13 def update_weight(w, b, dw, db, lr):
14     new_w = w - lr * dw
15     new_b = b - lr * db
16
17     return new_w, new_b
```

Code Implementation

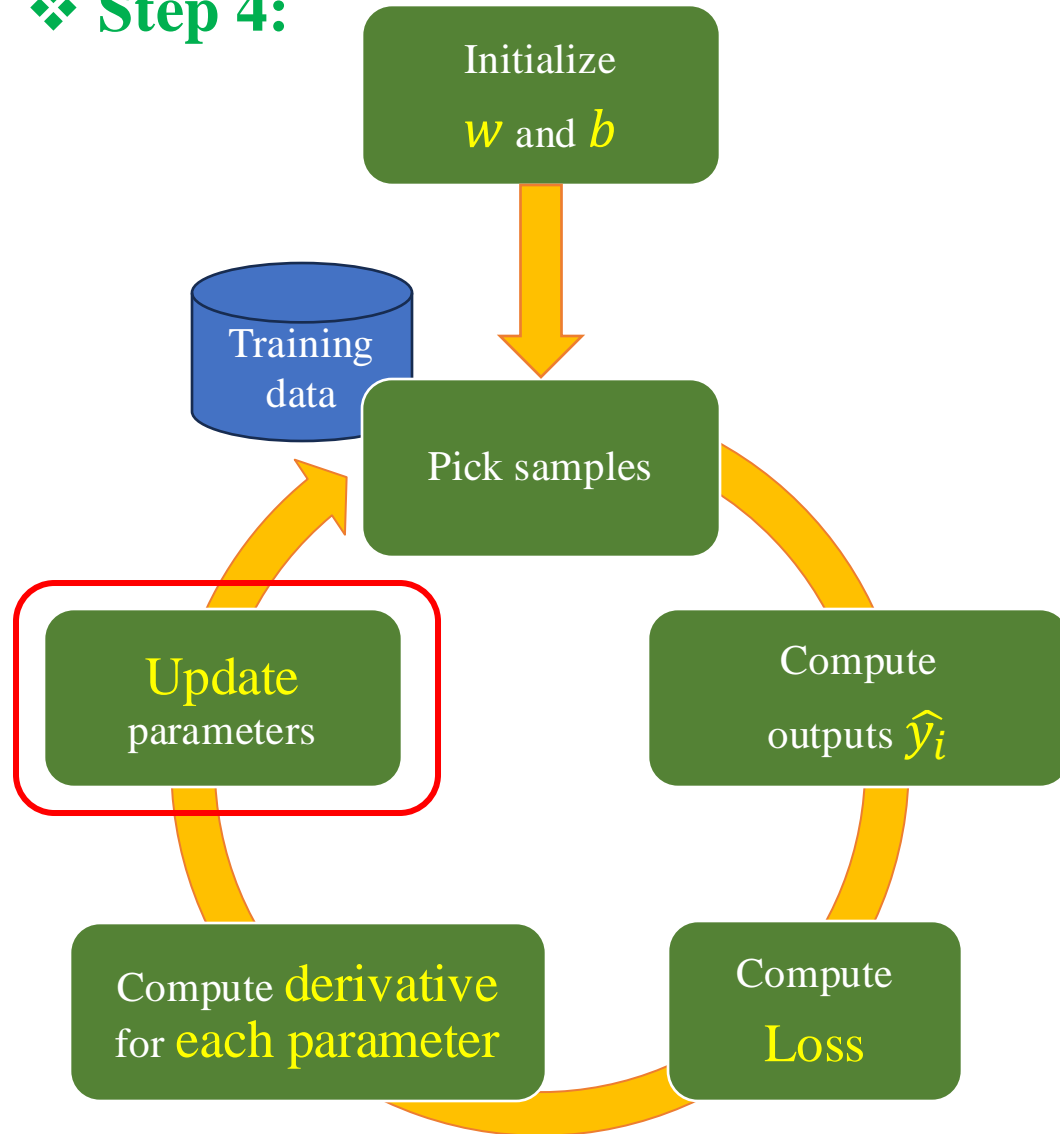
❖ Step 3:



```
1 def predict(x, w, b):  
2     return x * w + b  
3  
4 def compute_loss(y_hat, y):  
5     return (y_hat - y) ** 2  
6  
7 def compute_gradient(y_hat, y, x):  
8     dw = 2 * x * (y_hat - y)  
9     db = 2 * (y_hat - y)  
10  
11    return dw, db  
12  
13 def update_weight(w, b, dw, db, lr):  
14     new_w = w - lr * dw  
15     new_b = b - lr * db  
16  
17    return new_w, new_b
```


Code Implementation

❖ Step 4:



```
1 def predict(x, w, b):  
2     return x * w + b  
3  
4 def compute_loss(y_hat, y):  
5     return (y_hat - y) ** 2  
6  
7 def compute_gradient(y_hat, y, x):  
8     dw = 2 * x * (y_hat - y)  
9     db = 2 * (y_hat - y)  
10  
11     return dw, db  
12  
13 def update_weight(w, b, dw, db, lr):  
14     new_w = w - lr * dw  
15     new_b = b - lr * db  
16  
17     return new_w, new_b
```

Code Implementation

❖ Test functions

```
1 x = 5
2 y = 66
3
4 b = 5
5 w = 10
6 lr = 0.01
7
8 y_hat = predict(x, w, b)
9 print(f'Prediction: {y_hat}')
10
11 loss = compute_loss(y_hat, y)
12 print(f'Loss: {loss}')
13
14 dw, db = compute_gradient(y_hat, y, x)
15 print(f'Gradient: dw = {dw}, db = {db}')
16
17 new_w, new_b = update_weight(w, b, dw, db, lr)
18 print(f'Updated weight: w = {new_w}, b = {new_b}')
```

Prediction: 55

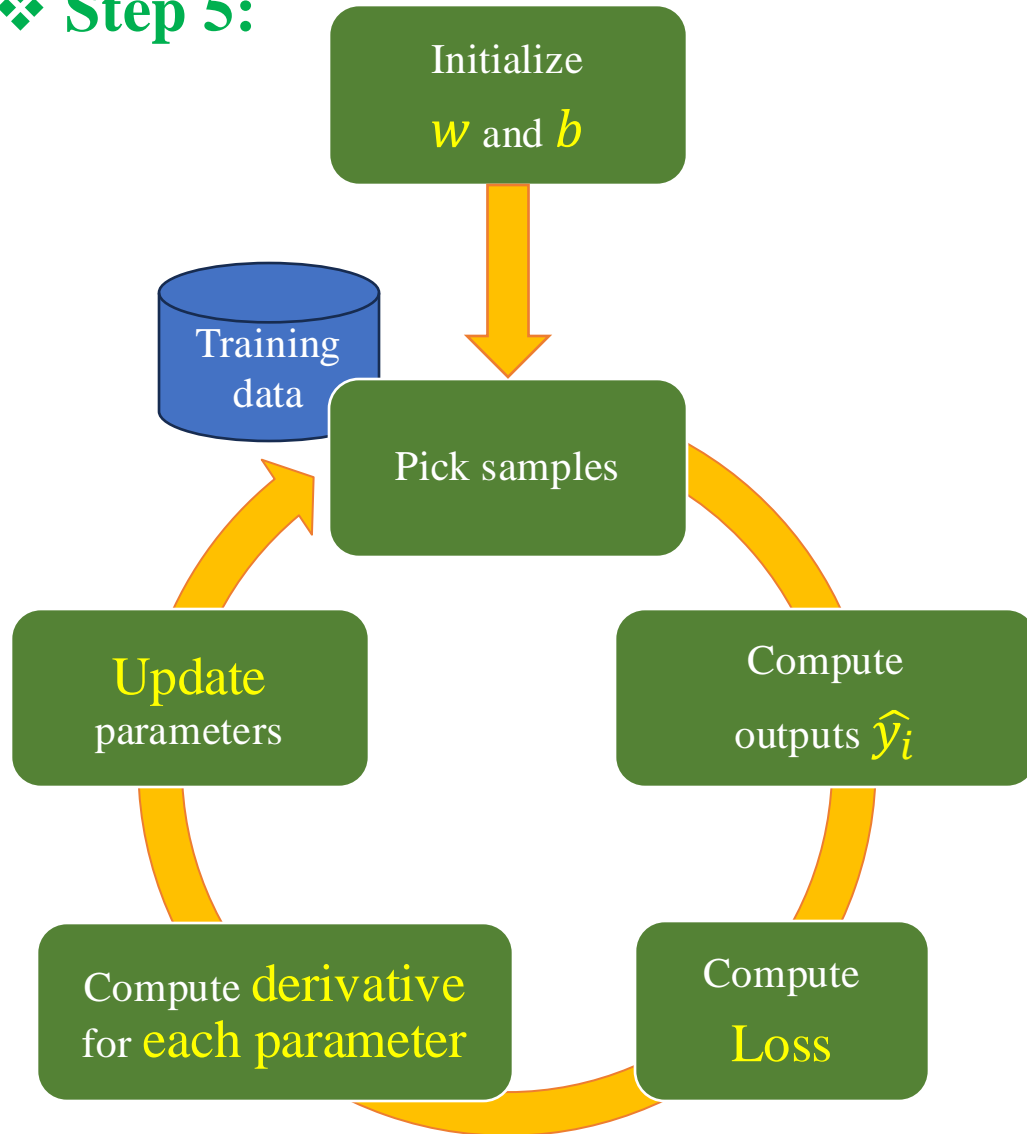
Loss: 121

Gradient: dw = -110, db = -22

Updated weight: w = 11.1, b = 5.22

Code Implementation

❖ Step 5:



```
1 X = [3, 4, 5, 6]
2 y = [60, 55, 66, 93]
3
4 epochs = 5
5
6 b = 5
7 w = 10
8 lr = 0.001
9
10 for epoch in range(epochs):
11     for idx in range(len(X)):
12         x = X[idx]
13         y_hat = predict(x, w, b)
14
15         loss = compute_loss(y_hat, y[idx])
16         dw, db = compute_gradient(y_hat, y[idx], x)
17         w, b = update_weight(w, b, dw, db, lr)
18
19     print(f'Epoch: {epoch}, Loss: {loss}, w: {w}, b: {b}')
```

Code Implementation

❖ Test functions

```
1 X = [3, 4, 5, 6]
2 y = [60, 55, 66, 93]
3
4 epochs = 5
5
6 b = 5
7 w = 10
8 lr = 0.001
9
10 for epoch in range(epochs):
11     for idx in range(len(X)):
12         x = X[idx]
13         y_hat = predict(x, w, b)
14
15         loss = compute_loss(y_hat, y[idx])
16         dw, db = compute_gradient(y_hat, y[idx], x)
17         w, b = update_weight(w, b, dw, db, lr)
18
19     print(f'Epoch: {epoch}, Loss: {loss}, w: {w}, b: {b}')
```

```
Epoch: 0, Loss: 625, w: 10.15, b: 5.05
Epoch: 0, Loss: 87.42250000000003, w: 10.2248, b: 5.0687
Epoch: 0, Loss: 96.18313328999996, w: 10.322873, b: 5.0883145999999995
Epoch: 0, Loss: 674.671917735367, w: 10.6345663688, b: 5.140263494799999
Epoch: 1, Loss: 526.9796530551046, w: 10.7723025931928, b: 5.186175569597599
Epoch: 1, Loss: 45.22043422409116, w: 10.826099505653849, b: 5.199624797712861
Epoch: 1, Loss: 44.48726818636239, w: 10.892798282394027, b: 5.212964553060897
Epoch: 1, Loss: 503.1159245209058, w: 11.161961231424927, b: 5.257825044566047
Epoch: 2, Loss: 451.8299181792317, w: 11.289498978991881, b: 5.300337627088366
Epoch: 2, Loss: 20.626734206131236, w: 11.325832310647435, b: 5.309420960002254
Epoch: 2, Loss: 16.495112001764532, w: 11.366446485515041, b: 5.317543794975775
Epoch: 2, Loss: 379.617577561682, w: 11.600251813018248, b: 5.356511349559643
Epoch: 3, Loss: 393.73406129822536, w: 11.719308212286561, b: 5.396196815982414
Epoch: 3, Loss: 7.434185791000412, w: 11.741120774965532, b: 5.4016499566521565
Epoch: 3, Loss: 3.5824880584478613, w: 11.760048236650734, b: 5.405435448989197
Epoch: 3, Loss: 290.1665292422299, w: 11.96445953822401, b: 5.43950399925141
Epoch: 4, Loss: 348.46127150556185, w: 12.07646224254047, b: 5.476838234023563
Epoch: 4, Loss: 1.4818504428538535, w: 12.086200744906987, b: 5.479272859615192
Epoch: 4, Loss: 0.008050291351769556, w: 12.087097979065486, b: 5.479452306446892
Epoch: 4, Loss: 224.93879873714388, w: 12.267073496895408, b: 5.509448226085213
```

Code Implementation

❖ Before and after training

Before training:

```
1 x = 5
2 y = 66
3
4 y_hat = predict(x, w, b)
5 print(f'Prediction: {y_hat}')
6
7 loss = compute_loss(y_hat, y)
8 print(f'Loss: {loss}')
9
10 print(f'w: {w}, b: {b}')
```

Prediction: 55
Loss: 121
w: 10, b: 5

After training:


```
1 x = 5
2 y = 66
3
4 y_hat = predict(x, w, b)
5 print(f'Prediction: {y_hat}')
6
7 loss = compute_loss(y_hat, y)
8 print(f'Loss: {loss}')
9
10 print(f'w: {w}, b: {b}')
```

Prediction: 66.84481571056226
Loss: 0.7137135848128137
w: 12.267073496895408, b: 5.509448226085213

Summarization and QA

Summarization

❖ Summarization



In this lecture, we have discussed:

1. Introduction to Linear Regression.
2. Simple
3. Use SVM to solve a classification and a regression tasks.

Question

