



AI VIET NAM

@aivietnam.edu.vn

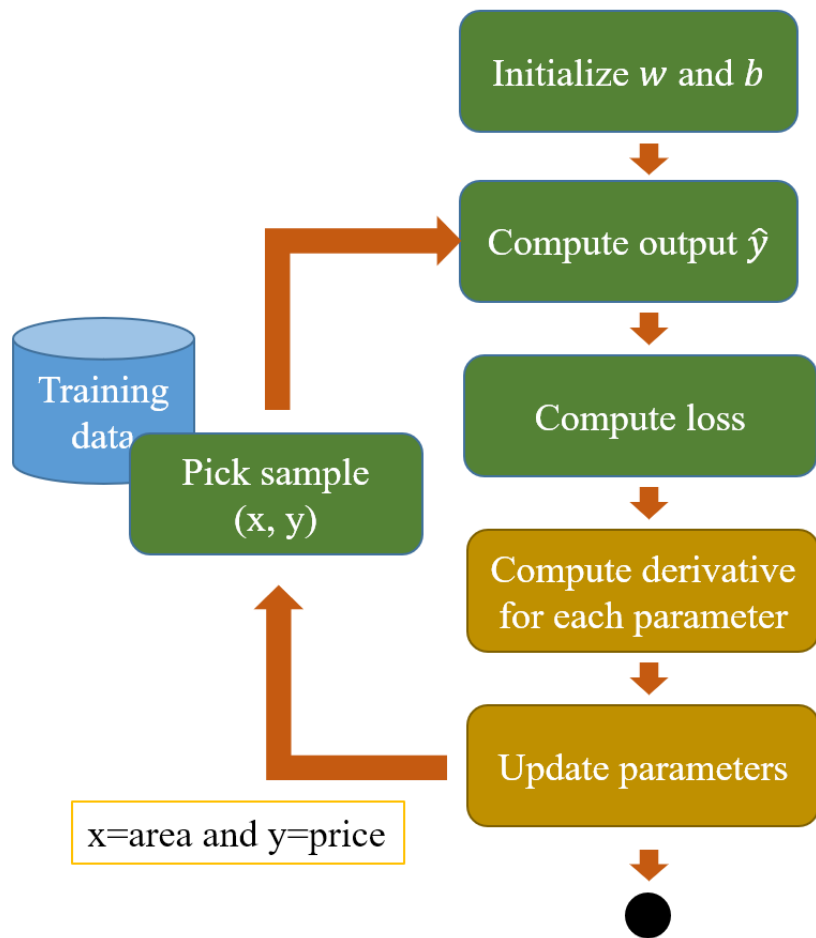
Linear Regression

A Simple Approach


Quang-Vinh Dinh
PhD in Computer Science

Objectives

Linear Regression



Computational Graph



Batch Training

1) Pick all the N samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < N$$

3) Compute loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < N$$

4) Compute derivatives

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$
$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)}) \quad \text{for } 0 \leq i < N$$

5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{N} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{N}$$

Outline

SECTION 1

Linear Regression

SECTION 2

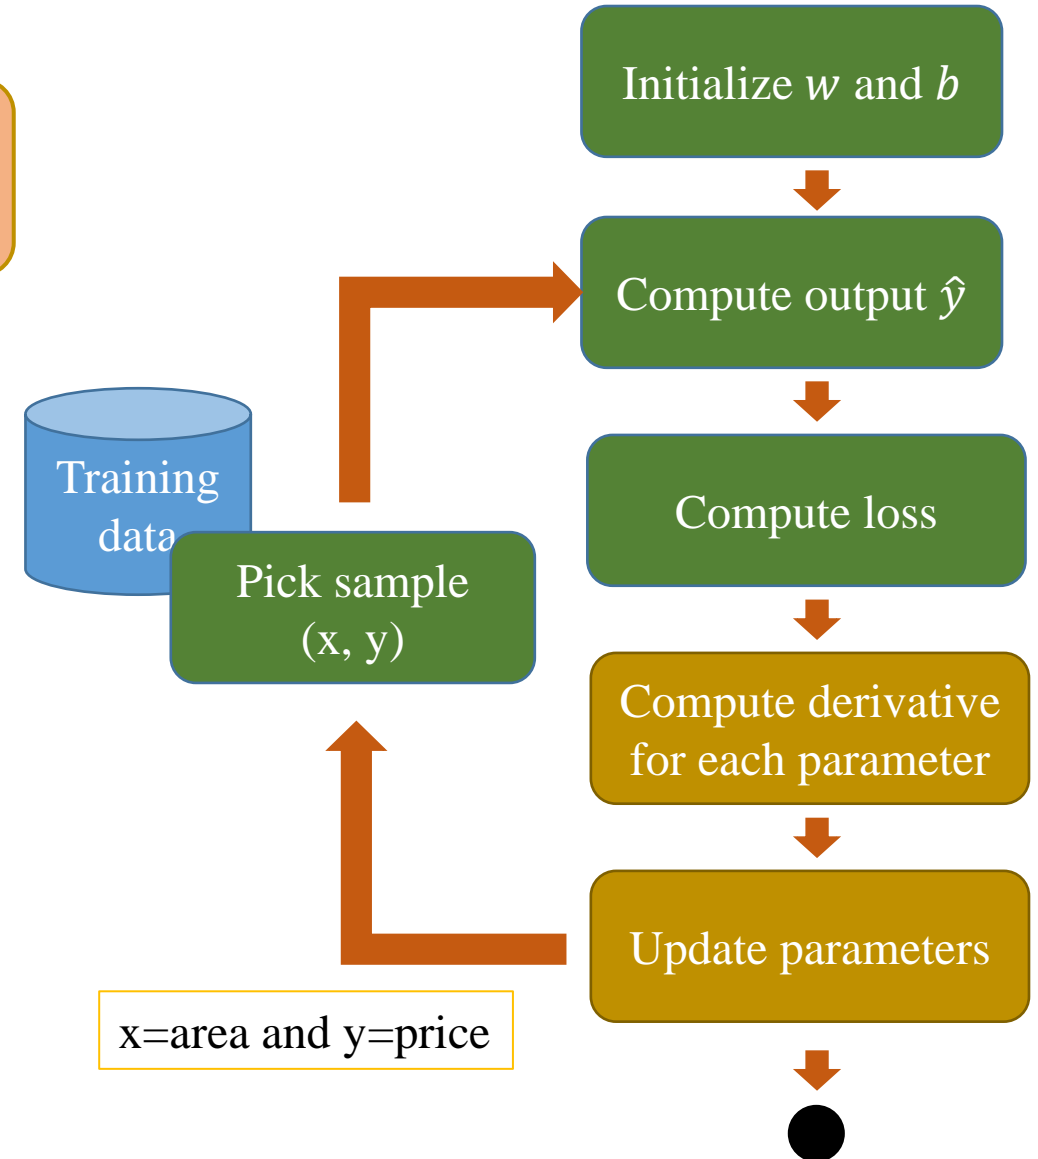
Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions



Linear Regression

Introduction

Feature		Label	
	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

House price data

Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising data

if area=6.0, price=?

if TV=55.0, Radio=34.0,
and Newspaper=62.0,
price=?


Features													Label
crim	↕ zn	↕ indus	↕ chas	↕ nox	↕ rm	↕ age	↕ dis	↕ rad	↕ tax	↕ ptratio	↕ black	↕ lstat	↕ medv
0.00632	18	2.31	0	0.538	6.575	65.2	4.09	1	296	15.3	396.9	4.98	24
0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.9	9.14	21.6
0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4
0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.9	5.33	36.2
0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.6	12.43	22.9

Boston House Price Data

House Price Prediction

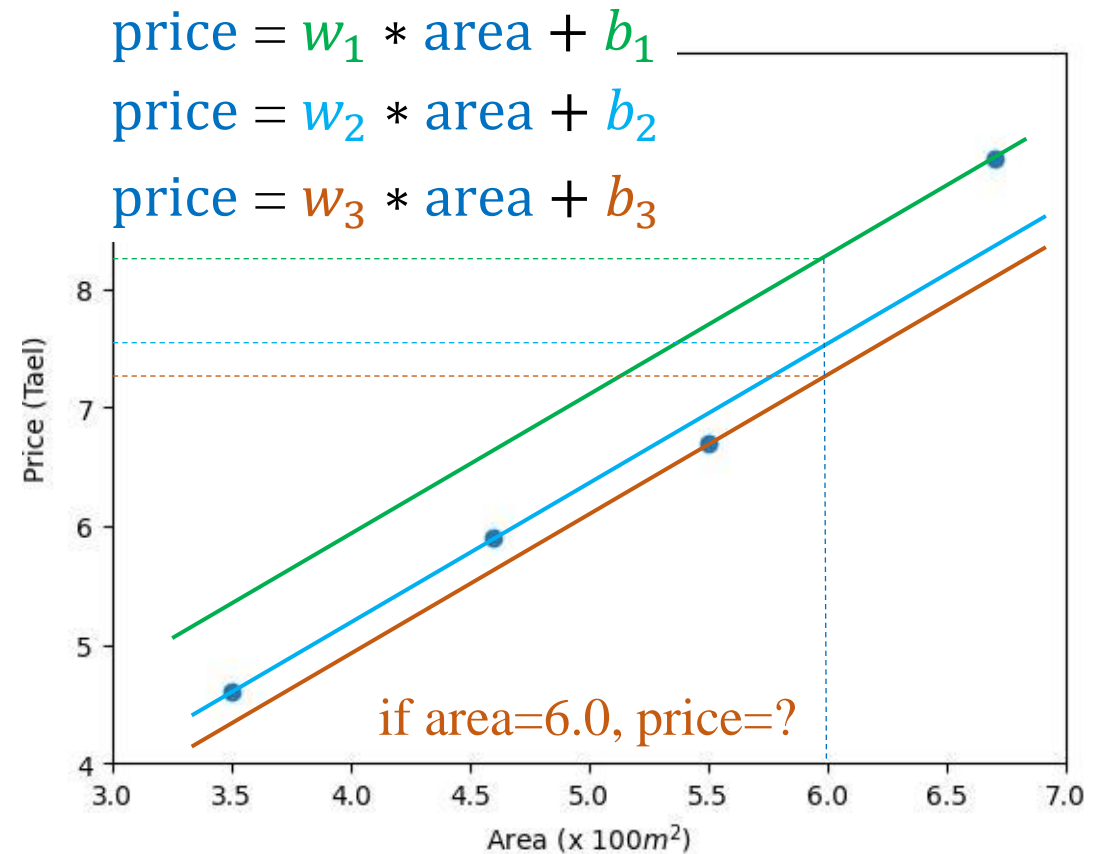
House price data

Feature	Label
area	price
6.7	8.1
4.6	5.6
3.5	4.3
5.5	6.7



Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7

House price data






Linear Regression

3

❖ Area-based house price prediction



$$error_i = distance(\hat{y}_i, y_i)$$



Find w and b whose model has the smallest error, where $error = \frac{1}{N} \sum_i error_i$ **How?**

❖ Area-based house price prediction

weight

bias

predicted_price = w * area + b


error = (predicted_price - real_price)²

$\hat{y}_i = wx_i + b$

$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$

area	price	predicted	error
6.7	9.1	-2.238	128.55
4.6	5.9	-1.524	55.11
3.5	4.6	-1.15	33.06
5.5	6.7	-1.83	72.76

$w = -0.34$
 $b = 0.04$



❖ Area-based house price prediction

$$\text{predicted_price} = w * \text{area} + b$$

$$\text{error} = (\text{predicted_price} - \text{real_price})^2$$


$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

area	price	predicted	error
6.7	9.1	8.099	1.002
4.6	5.9	5.642	0.066
3.5	4.6	4.355	0.06
5.5	6.7	6.695	0.00002

$$w = 1.17$$



$$b = 0.26$$



❖ Area-based house price prediction

$$\hat{y}_i = wx_i + b$$
$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

How to change w and b
so that $L(\hat{y}_i, y_i)$ reduces




Linear Regression

❖ Understanding the loss function


$$\hat{y}_i = wx_i + b$$

$$L(\hat{y}_i, y_i) = (\hat{y}_i - y_i)^2$$

How to change w and b so that $L(\hat{y}_i, y_i)$ reduces



Different b values with a fixed w value



Different w values with a fixed b value

Linear Regression

Linear equation

$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,

w and b are parameters

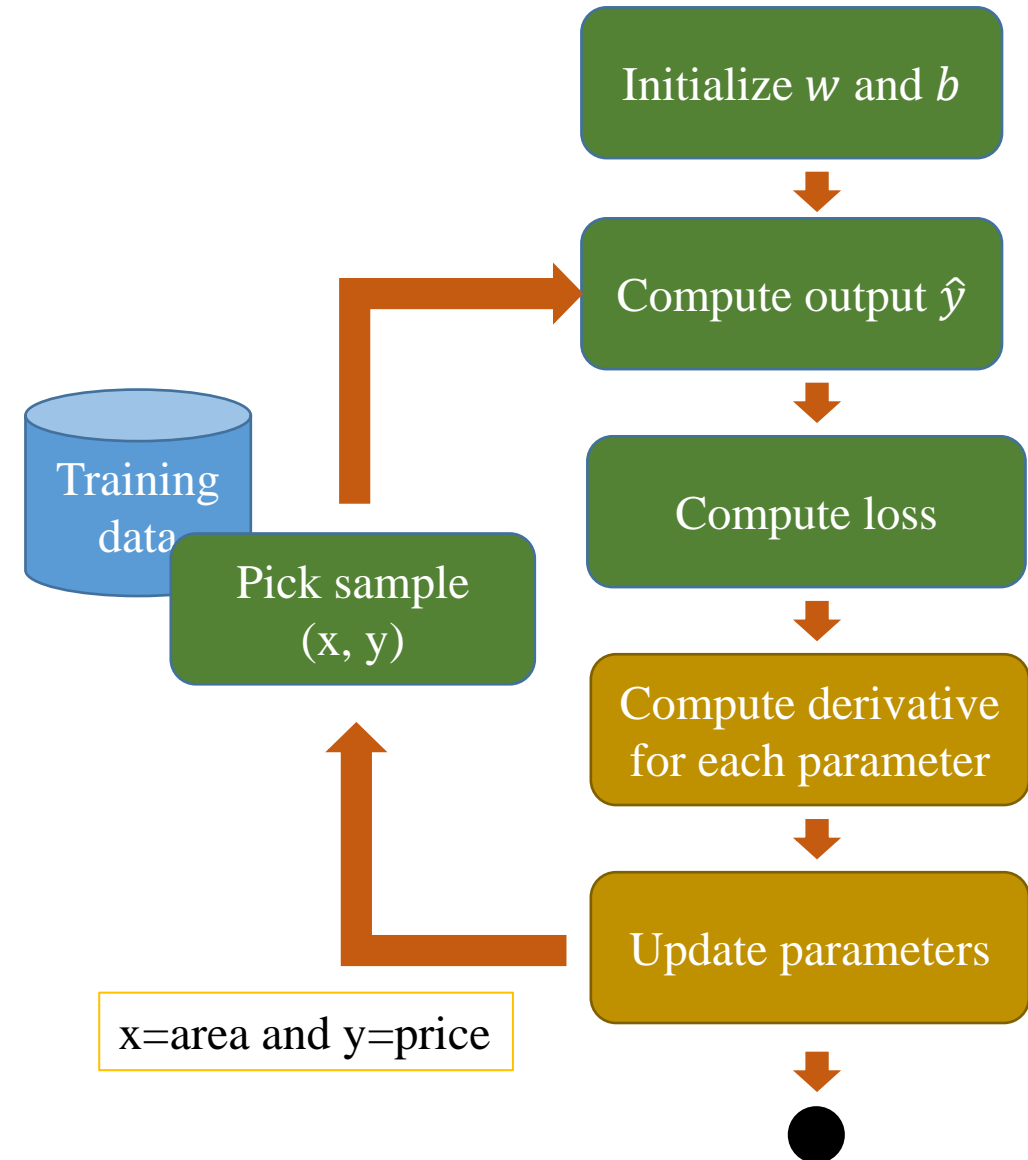
and x is an input feature

Error (loss) computation

Idea: compare predicted values \hat{y} and label values y

Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$



Linear equation

$$\hat{y} = wx + b$$

where \hat{y} is a predicted value,

w and b are parameters

and x is an input feature

Error (loss) computation

Idea: compare predicted values \hat{y} and label values y

Squared loss

$$L(\hat{y}, y) = (\hat{y} - y)^2$$

Find better w and b

Use gradient descent to minimize the loss function

Compute derivate for each parameter

$$\frac{\partial L}{\partial w} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = \frac{\partial L}{\partial \hat{y}} \frac{\partial \hat{y}}{\partial b} = 2(\hat{y} - y)$$

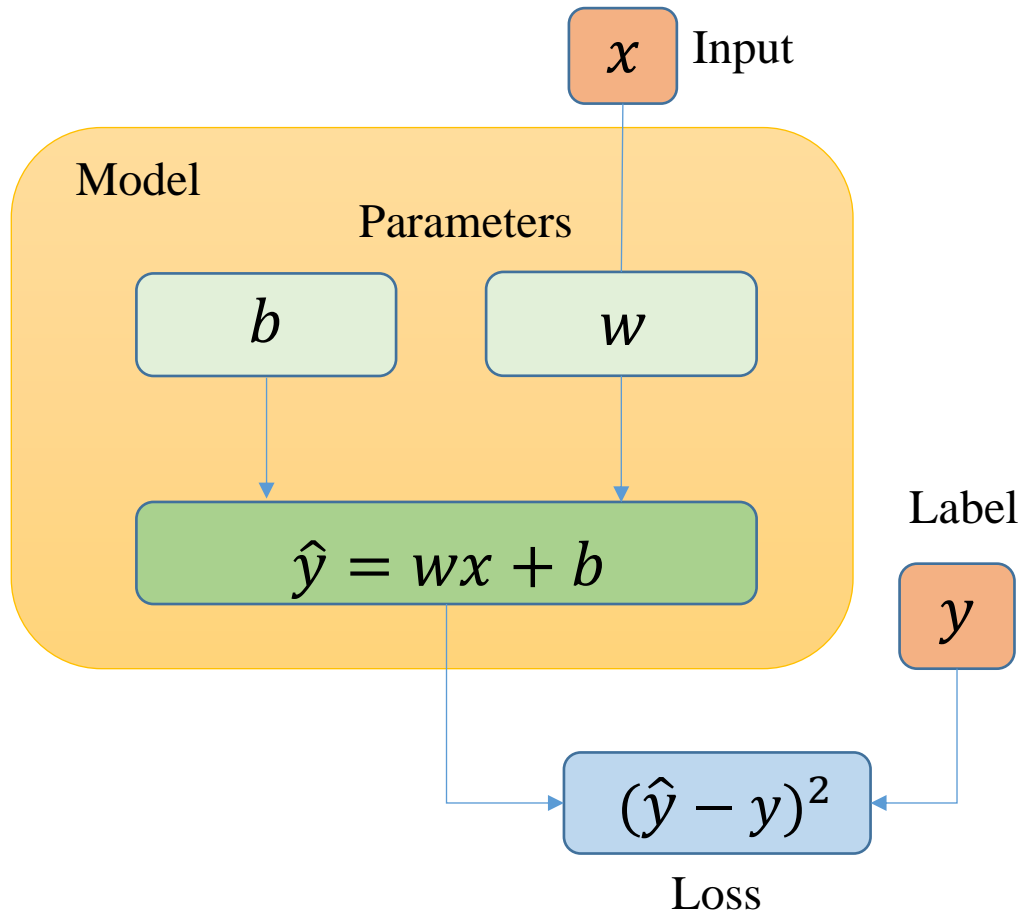
Update parameters

$$w = w - \eta \frac{\partial L}{\partial w} \quad b = b - \eta \frac{\partial L}{\partial b}$$

η is learning rate

❖ Simple example

Diagram



Cheat sheet

Compute the output \hat{y}

$$\hat{y} = wx + b$$

Compute the loss

$$L = (\hat{y} - y)^2$$

Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$


$$b = b - \eta \frac{\partial L}{\partial b}$$

Linear Regression

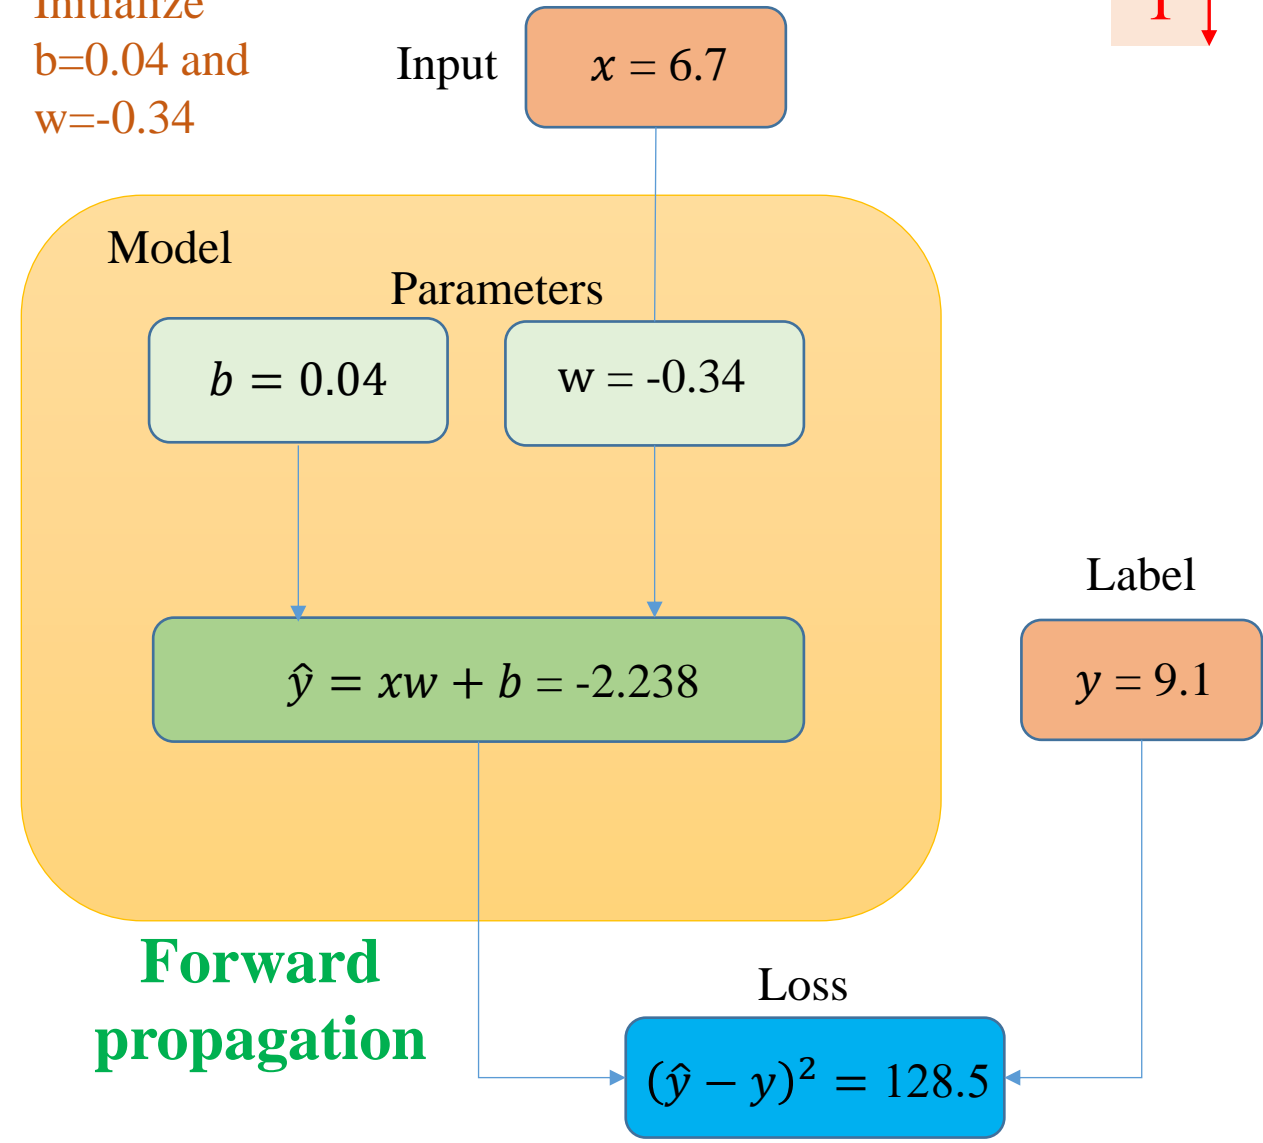
11

Given
sample
data

Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7



Initialize
 $b=0.04$ and
 $w=-0.34$



Linear Regression

12

2

Input

$x = 6.7$

Backpropagation

$\eta = 0.01$

Model

Parameters

$b = 0.26676$

$w = 1.17929$

$$b = b - \eta \frac{\partial L}{\partial b}$$

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$\hat{y} = xw + b = -2.238$$

$$\begin{aligned} \frac{\partial L}{\partial w} &= 2x(\hat{y} - y) \\ &= -151.9292 \end{aligned}$$

$$\begin{aligned} \frac{\partial L}{\partial b} &= 2(\hat{y} - y) \\ &= -22.676 \end{aligned}$$

Label

$y = 9.1$

Loss

$$(\hat{y} - y)^2 = 128.5$$

3

Input

$x = 6.7$

Forward propagation

Model

Parameters

$b = 0.26676$

$w = 1.17929$

$$b = b - \eta \frac{\partial L}{\partial b}$$

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$\hat{y} = xw + b = 8.168$$

Label

$y = 9.1$


Loss

$$(\hat{y} - y)^2 = 0.868$$

New w and b help
the loss reduce


❖ Simple example

Model prediction before and after the first update



$w = -0.34$ $b = 0.04$ $L = 128.55$

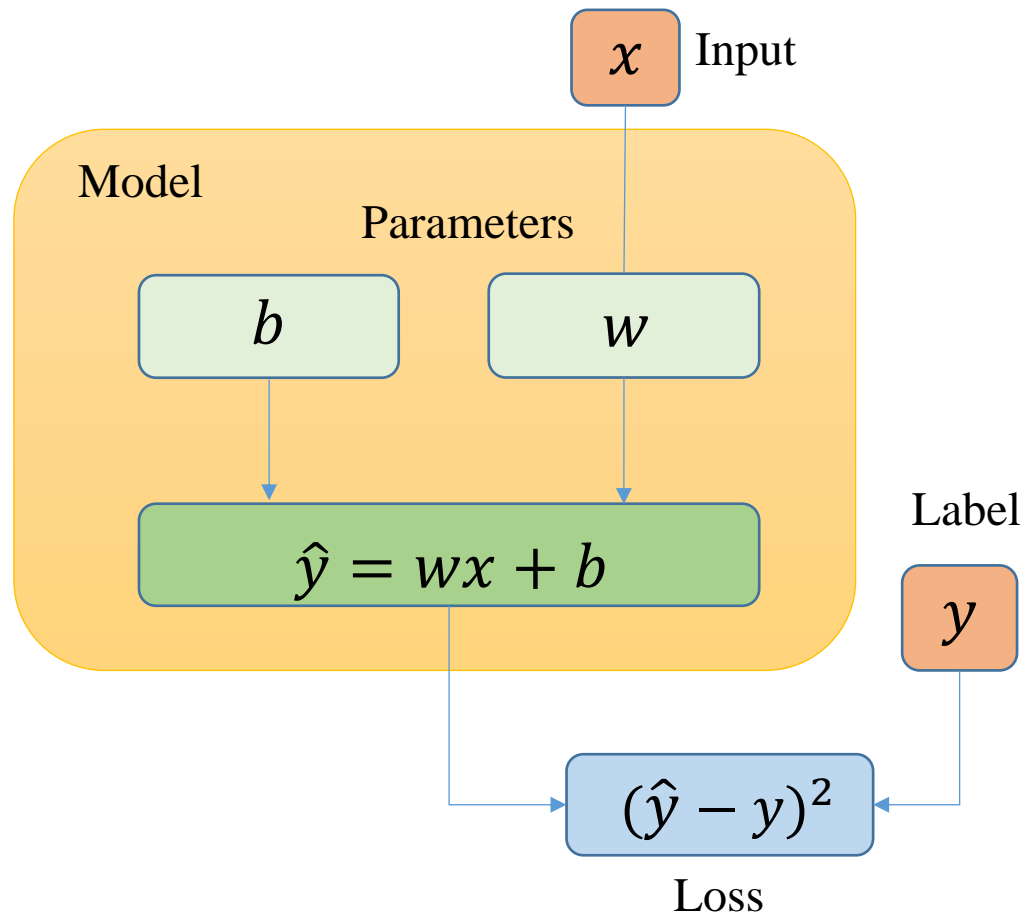
Before updating



$w = 1.179292$ $b = 0.26676$ $L = 0.868$

After updating

❖ Summary (simple version)



1) Pick a sample (x, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = wx + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

η is learning rate

❖ Implementation

Cheat sheet

Compute the output \hat{y}

$$\hat{y} = wx + b$$

Compute the loss

$$L = (\hat{y} - y)^2$$

Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y)$$

$$\frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

Update parameters

$$w = w - \eta \frac{\partial L}{\partial w}$$

$$b = b - \eta \frac{\partial L}{\partial b}$$

```
1 # forward
2 def predict(x, w, b):
3     return x*w + b
4
5 # compute gradient
6 def gradient(y_hat, y, x):
7     dw = 2*x*(y_hat-y)
8     db = 2*(y_hat-y)
9
10    return (dw, db)
11
12 # update weights
13 def update_weight(w, b, lr, dw, db):
14     w_new = w - lr*dw
15     b_new = b - lr*db
16
17    return (w_new, b_new)
```


Computational Graph (A different viewpoint)

❖ House price prediction

❖ One-sample training


$$price = w * area + b$$

$$t = w * area$$



❖ House price prediction

❖ One-sample training




w and b are usually set by random values

For example, $N(0, \sigma)$ where σ is small

❖ House price prediction


❖ One-sample training



Feature		Label	
area		price	
6.7	9.1		
4.6	5.9		
3.5	4.6		
5.5	6.7		


❖ House price prediction

❖ One-sample training



❖ House price prediction

❖ One-sample training



Update w and b

$$w = w - \eta * \frac{\partial L}{\partial w}$$

$$b = b - \eta * \frac{\partial L}{\partial b}$$


Learning rate $\eta = 0.01$

$$w = -0.34 - (0.01 * (-151.9)) = 1.179$$

$$b = 0.04 - (0.01 * (-22.67)) = 0.266$$

❖ House price prediction

❖ One-sample training




	Feature	Label	
	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

❖ House price prediction

❖ One-sample training

	Feature	Label
	area	price
	6.7	9.1
	4.6	5.9
	3.5	4.6
	5.5	6.7



Updated a and b values help to reduce the L value

Implementation

❖ One-sample training

Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7

column index=0
column index=1

```
1 # data preparation
2 import numpy as np
3 import matplotlib.pyplot as plt
4
5 def get_column(data, index):
6     result = [row[index] for row in data]
7     return result
8
9 data = np.genfromtxt('data.csv',
10                      delimiter=',').tolist()
11
12 x_data = get_column(data, 0)
13 y_data = get_column(data, 1)
14 N = len(x_data)
15
16 print(f'areas: {x_data}')
17 print(f'prices: {y_data}')
18 print(f'data_size: {N}')
```

```
areas: [6.7, 4.6, 3.5, 5.5]
prices: [9.1, 5.9, 4.6, 6.7]
data_size: 4
```

❖ House price prediction

❖ One-sample training

1) Pick a sample (x, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = wx + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w} = 2x(\hat{y} - y) \quad \frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

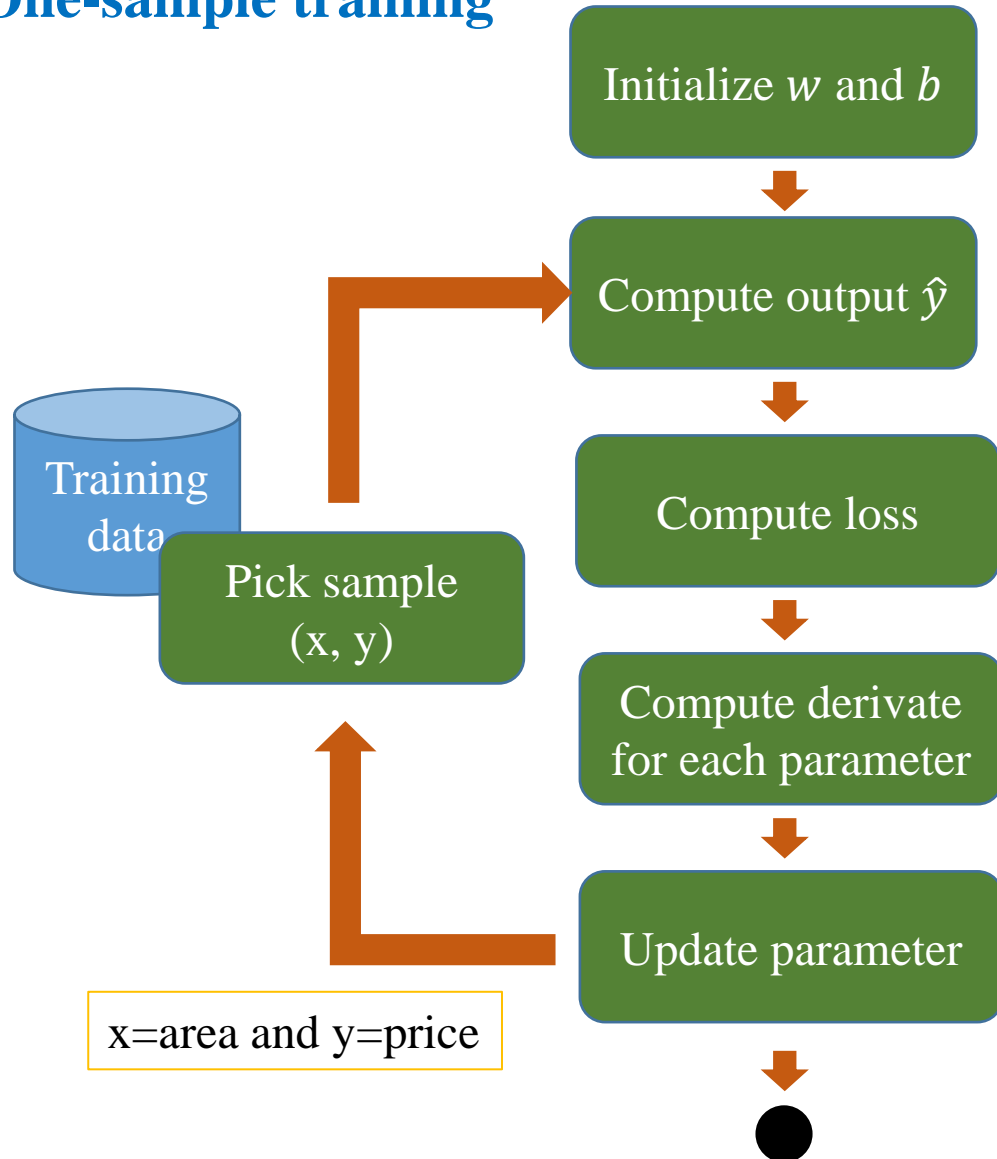
5) Update parameters

$$w = w - \eta \frac{\partial L}{\partial w} \quad b = b - \eta \frac{\partial L}{\partial b}$$

```
1 # forward
2 def predict(x, w, b):
3     return x*w + b
4
5 # compute gradient
6 def gradient(y_hat, y, x):
7     dw = 2*x*(y_hat-y)
8     db = 2*(y_hat-y)
9
10    return (dw, db)
11
12 # update weights
13 def update_weight(w, b, lr, dw, db):
14     w_new = w - lr*dw
15     b_new = b - lr*db
16
17    return (w_new, b_new)
```

Implementation

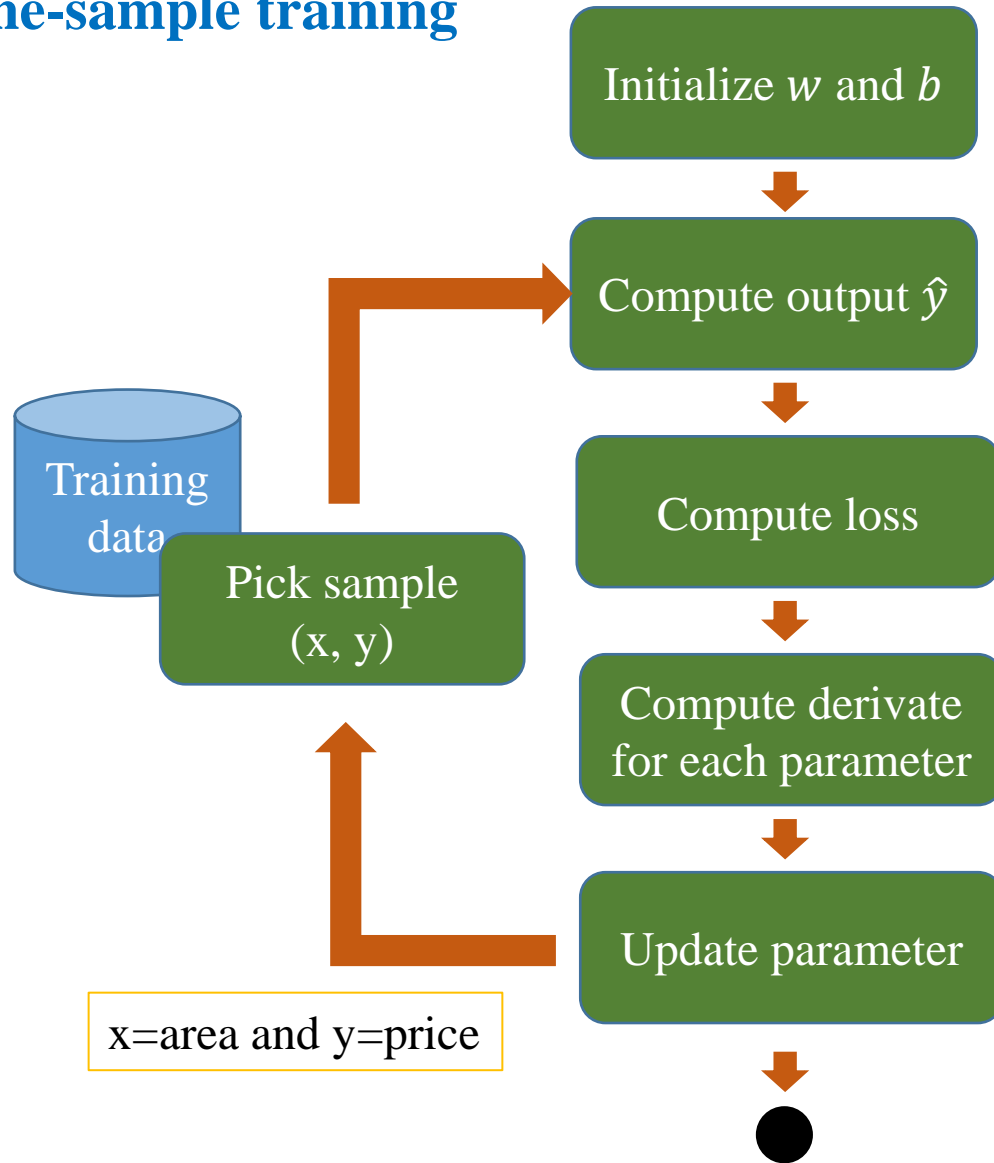
❖ One-sample training



```
1 # init weights
2 b = 0.04
3 w = -0.34
4 lr = 0.01
5
6 # how long
7 epoch_max = 10
8
9 for epoch in range(epoch_max):
10     for i in range(data_size):
11         # get a sample
12         # ...
13
14         # predict y_hat
15         # ...
16
17         # compute loss
18         # ...
19
20         # compute gradient
21         # ...
22
23         # update weights
24         # ...
25
```

Implementation


❖ One-sample training





```
1 # init weights
2 b = 0.04
3 w = -0.34
4 lr = 0.01
5
6 # how long
7 epoch_max = 10
8 data_size = 4
9
10 for epoch in range(epoch_max):
11     for i in range(data_size):
12         # get a sample
13         x = areas[i]
14         y = prices[i]
15
16         # predict y_hat
17         y_hat = predict(x, w, b)
18
19         # compute loss
20         loss = (y_hat-y)*(y_hat-y)
21
22         # compute gradient
23         (dw, db) = gradient(y_hat, y, x)
24
25         # update weights
26         (w, b) = update_weight(w, b, lr, dw, db)
```

Implementation

- ❖ House price prediction
 - ❖ One-sample training



Model training



Quiz 1: Construct for the following data

Features		Label
↕ Radio	↕ Newspaper	↕ Sales
37.8	69.2	22.1
39.3	45.1	10.4
45.9	69.3	12
41.3	58.5	16.5
10.8	58.4	17.9

Outline

SECTION 1

Linear Regression

SECTION 2

Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions

1) Pick m samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < m$$

3) Compute loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < m$$


4) Compute derivatives

$$\begin{aligned} \frac{\partial L^{(i)}}{\partial w} &= 2x^{(i)}(\hat{y}^{(i)} - y^{(i)}) \\ \frac{\partial L^{(i)}}{\partial b} &= 2(\hat{y}^{(i)} - y^{(i)}) \end{aligned} \quad \text{for } 0 \leq i < m$$

5) Update


$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{m} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{m}$$

❖ Compute derivate for w and b




mini-batch $m = 2$

❖ Compute derivate for w and b




mini-batch $m = 2$

❖ Compute derivate for w and b




❖ Compute derivate for w and b



$$\frac{\partial L}{\partial w} = \frac{1}{m} \sum_i 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$

❖ Compute derivate for w and b

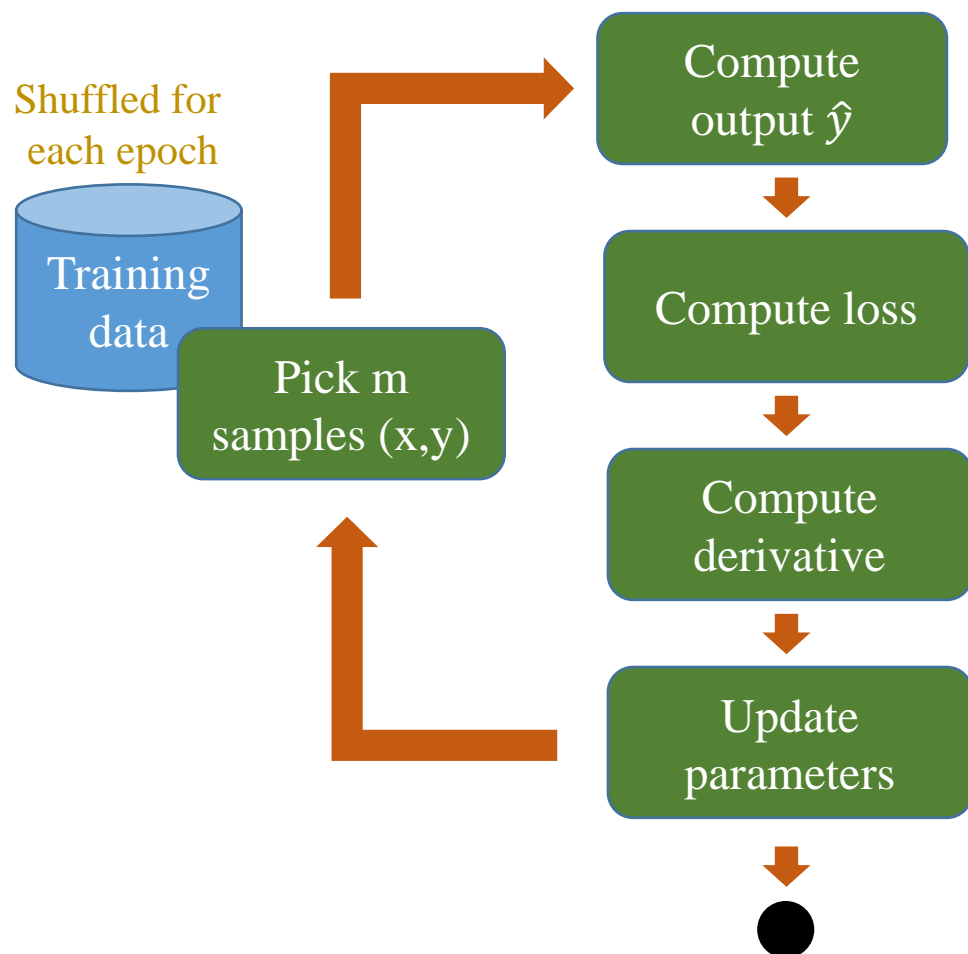


$$\frac{\partial L}{\partial w} = \frac{1}{m} \sum_i 2x^{(i)} (\hat{y}^{(i)} - y^{(i)})$$

$$\frac{\partial L}{\partial b} = \frac{1}{m} \sum_i 2(\hat{y}^{(i)} - y^{(i)})$$

❖ House price prediction

❖ m-sample training ($1 < m < N$)



1) Pick m samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < m$$

3) Compute loss

$$L = \frac{1}{m} \sum_i (\hat{y}^{(i)} - y^{(i)})^2$$

4) Compute derivatives

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$

$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)})$$


for $0 \leq i < m$

5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{m} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{m}$$

❖ House price prediction

❖ m-sample training ($1 < m < N$)




Computational graph

❖ House price prediction

❖ m-sample training ($1 < m < N$)

Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7

$m = 2$



Computational graph

❖ House price prediction

❖ m-sample training ($1 < m < N$)

Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7

$m = 2$

Update w and b

$$w = w - \eta * \frac{\partial L}{\partial w}$$

$$b = b - \eta * \frac{\partial L}{\partial b}$$

Learning rate $\eta = 0.01$

replace

replace

$$x = \begin{bmatrix} 6.7 \\ 4.6 \end{bmatrix}$$

$$w = -0.34$$

$$\frac{\partial L}{\partial w} = -110.115$$

$$b = 0.04$$

$$\frac{\partial L}{\partial b} = -18.762$$

$$y = \begin{bmatrix} 9.1 \\ 5.9 \end{bmatrix}$$

*

t

+

$(\cdot)^2$

$$t = wx$$

$$\hat{y} = t + b$$


$$L = (\hat{y} - y)^2$$

$$w = -0.34 - (0.01 * (-110.115)) = 0.761$$


$$b = 0.04 - (0.01 * (-18.762)) = 0.227$$

❖ House price prediction


❖ m-sample training ($1 < m < N$)



Losses for 30 iterations



Model updating for different iterations



Outline

SECTION 1

Linear Regression

SECTION 2

Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions

1) Pick all the N samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < N$$

3) Compute loss

$$L^{(i)} = (\hat{y}^{(i)} - y^{(i)})^2 \quad \text{for } 0 \leq i < N$$

4) Compute derivatives

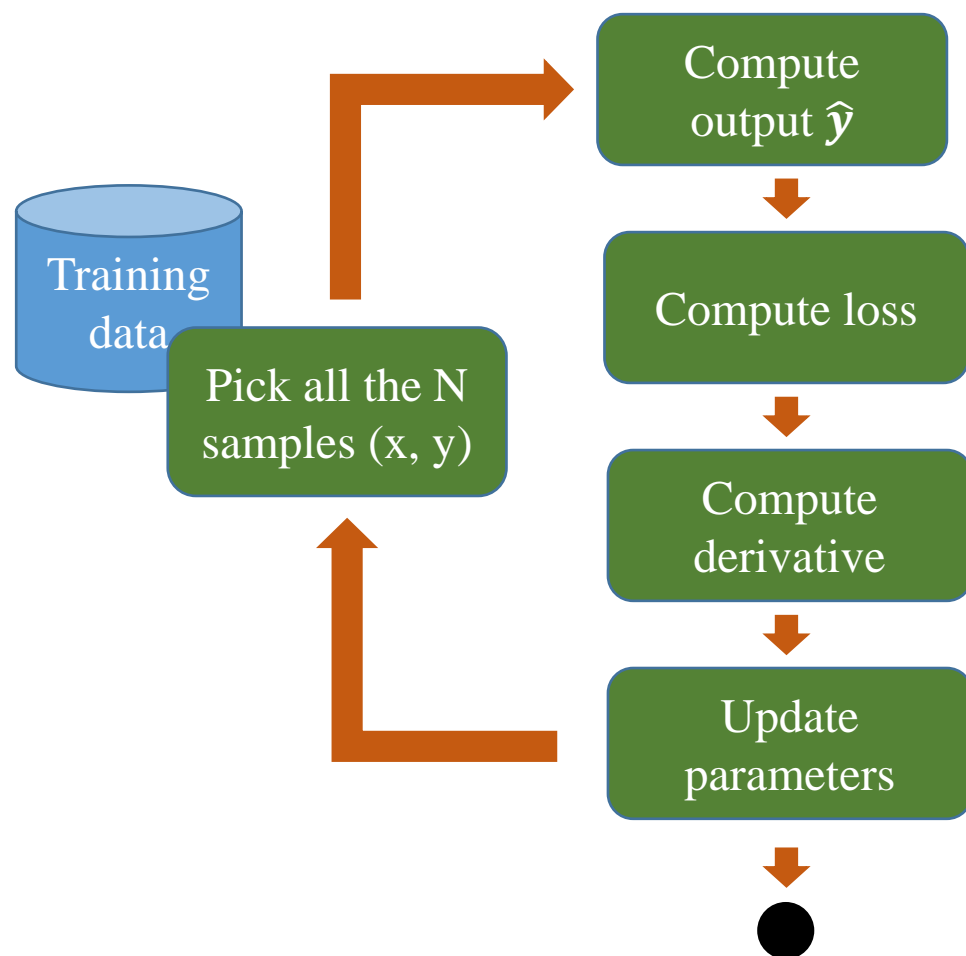
$$\begin{aligned} \frac{\partial L^{(i)}}{\partial w} &= 2x^{(i)}(\hat{y}^{(i)} - y^{(i)}) \\ \frac{\partial L^{(i)}}{\partial b} &= 2(\hat{y}^{(i)} - y^{(i)}) \end{aligned} \quad \text{for } 0 \leq i < N$$

5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{N} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{N}$$

❖ House price prediction

❖ N-sample training



1) Pick all the N samples $(x^{(i)}, y^{(i)})$ from training data

2) Compute output $\hat{y}^{(i)}$

$$\hat{y}^{(i)} = wx^{(i)} + b \quad \text{for } 0 \leq i < N$$

3) Compute loss

$$L = \frac{1}{N} \sum_i (\hat{y}^{(i)} - y^{(i)})^2$$

4) Compute derivatives

$$\frac{\partial L^{(i)}}{\partial w} = 2x^{(i)}(\hat{y}^{(i)} - y^{(i)})$$


$$\frac{\partial L^{(i)}}{\partial b} = 2(\hat{y}^{(i)} - y^{(i)}) \quad \text{for } 0 \leq i < N$$

5) Update

$$w = w - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial w}}{N} \quad b = b - \eta \frac{\sum_i \frac{\partial L^{(i)}}{\partial b}}{N}$$

❖ House price prediction

❖ N-sample training




Computational graph

❖ House price prediction

❖ N-sample training

$$\frac{\text{sum}(\frac{\partial L}{\partial w})}{4} = -88.5775$$

$$\frac{\text{sum}(\frac{\partial L}{\partial b})}{4} = -16.521$$



❖ House price prediction

❖ N-sample training

Update w and b


$$w = w - \eta * \frac{\partial L}{\partial w}$$

$$b = b - \eta * \frac{\partial L}{\partial b}$$

Learning rate $\eta = 0.01$


$$w = -0.34 - (0.01 * (-88.5775)) = 0.54$$

$$b = 0.04 - (0.01 * (-16.521)) = 0.205$$




❖ House price prediction

❖ N-sample training



Losses for 30 iterations



Model updating for different iterations

Extension

Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising data

Model: $\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b$

$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$

❖ General formula

	Feature	Label	
	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

House price data

Model: $\hat{y} = w_1x_1 + b$
 price = a * area + b

Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising data

Model: $\hat{y} = w_1x_1 + w_2x_2 + w_3x_3 + b$
 Sale = $w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$

Linear Regression

1) Pick a sample (x_1, x_2, x_3, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = w_1 * TV + w_2 * R + w_3 * N + b$$

$$\hat{y} = w_1 * x_1 + w_2 * x_2 + w_3 * x_3 + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w_1} = 2x_1(\hat{y} - y) \quad \frac{\partial L}{\partial w_3} = 2x_3(\hat{y} - y)$$

$$\frac{\partial L}{\partial w_2} = 2x_2(\hat{y} - y) \quad \frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w_1 = w_1 - \eta \frac{\partial L}{\partial w_1} \quad w_3 = w_3 - \eta \frac{\partial L}{\partial w_3}$$

$$w_2 = w_2 - \eta \frac{\partial L}{\partial w_2} \quad b = b - \eta \frac{\partial L}{\partial b}$$

Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising data

Model

$$\text{Sale} = w_1 * TV + w_2 * Radio + w_3 * Newspaper + b$$

$$\hat{y} = w_1 x_1 + w_2 x_2 + w_3 x_3 + b$$

1) Pick a sample (x_1, x_2, x_3, y) from training data

2) Compute the output \hat{y}

$$\hat{y} = w_1 * TV + w_2 * R + w_3 * N + b$$

$$\hat{y} = w_1 * x_1 + w_2 * x_2 + w_3 * x_3 + b$$

3) Compute loss

$$L = (\hat{y} - y)^2$$

4) Compute derivative

$$\frac{\partial L}{\partial w_1} = 2x_1(\hat{y} - y) \quad \frac{\partial L}{\partial w_3} = 2x_3(\hat{y} - y)$$

$$\frac{\partial L}{\partial w_2} = 2x_2(\hat{y} - y) \quad \frac{\partial L}{\partial b} = 2(\hat{y} - y)$$

5) Update parameters

$$w_1 = w_1 - \eta \frac{\partial L}{\partial w_1} \quad w_3 = w_3 - \eta \frac{\partial L}{\partial w_3}$$

$$w_2 = w_2 - \eta \frac{\partial L}{\partial w_2} \quad b = b - \eta \frac{\partial L}{\partial b}$$

```
1  # compute output and loss
2  def predict(x1, x2, x3, w1, w2, w3, b):
3      return w1*x1 + w2*x2 + w3*x3 + b
4  def compute_loss(y_hat, y):
5      return (y_hat - y)**2
6
7  # compute gradient
8  def compute_gradient_wi(xi, y, y_hat):
9      dl_dwi = 2*xi*(y_hat-y)
10     return dl_dwi
11 def compute_gradient_b(y, y_hat):
12     dl_db = 2*(y_hat-y)
13     return dl_db
14
15 # update weights
16 def update_weight_wi(wi, dl_dwi, lr):
17     wi = wi - lr*dl_dwi
18     return wi
19 def update_weight_b(b, dl_db, lr):
20     b = b - lr*dl_db
21     return b
```

Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

```

1 def initialize_params():
2     w1 = random.gauss(mu=0.0, sigma=0.01)
3     w2 = random.gauss(mu=0.0, sigma=0.01)
4     w3 = random.gauss(mu=0.0, sigma=0.01)
5     b = 0
6
7     return w1, w2, w3, b
8
9 # initialize model's parameters
10 w1, w2, w3, b = initialize_params()
11 print(w1, w2, w3, b)

```

0.01609506469549467 0.00607778501208891 0.0023344573891806507 0

```

1 import numpy as np
2 import random
3
4 def get_column(data, index):
5     result = [row[index] for row in data]
6     return result
7
8 data = np.genfromtxt('advertising.csv',
9                     delimiter=',',
10                    skip_header=1).tolist()
11
12 # get tv (index=0)
13 tv_data = get_column(data, 0)
14
15 # get radio (index=1)
16 radio_data = get_column(data, 1)
17
18 # get newspaper (index=2)
19 newspaper_data = get_column(data, 2)
20
21 # get sales (index=3)
22 sales_data = get_column(data, 3)

```

Outline

SECTION 1

Linear Regression

SECTION 2


Mini-batch Training

SECTION 3

Batch Training

SECTION 4

Loss Functions (Optional)



Discussion 1: Is it OK to use the following loss function?

$$L = \frac{1}{2}(\hat{y} - y)^2$$

Discussion 2: if so, construct formulas

Discussion 3: What about the following loss function?

$$L = \frac{1}{2} (y - \hat{y})^2$$


Discussion 4: Construct the connection between the two following losses?

$$L_1 = \frac{1}{2} (\hat{y} - y)^2$$

$$L_2 = (\hat{y} - y)^2$$

$$\hat{y} = wx + b$$

Discussion 5:
Can we remove b?



$$\hat{y} = wx + b$$

Discussion 5:
Can we remove b?

